



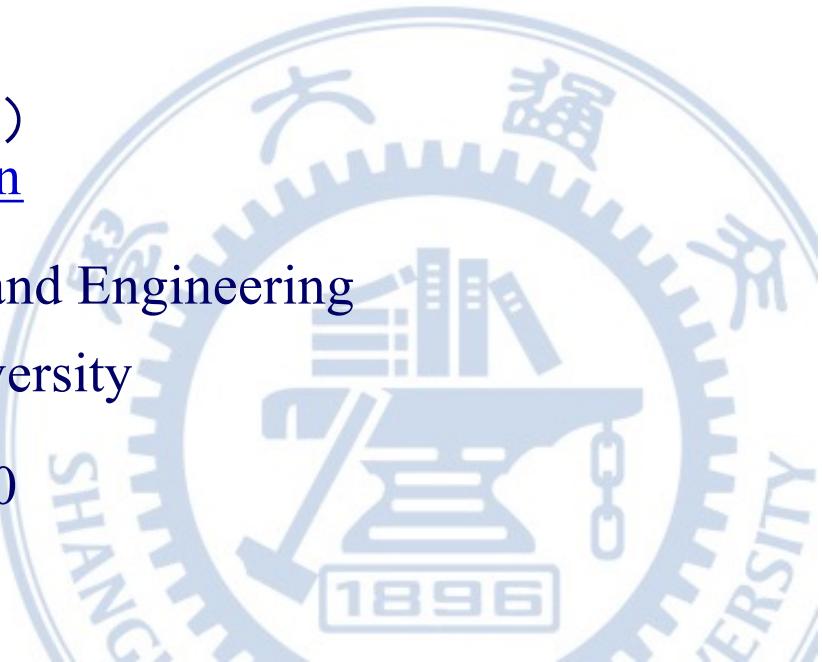
m.i.n Institute of Media,
Information, and Network

Deep Learning

Wenrui Dai (戴文睿)
<http://min.sjtu.edu.cn>

Department of Computer Science and Engineering
Shanghai Jiao Tong University

December 2nd, 2020



Today's Topics

- ◆ **Fundamentals of Deep Learning**
- ◆ **Convolutional Neural Networks (CNNs)**
- ◆ **Recurrent Neural Networks (RNNs)**
- ◆ **Generative Models**

Introduction

Forbes Billionaires Innovation Leadership Money Consumer Industry

Science Home News Journals Topics Careers

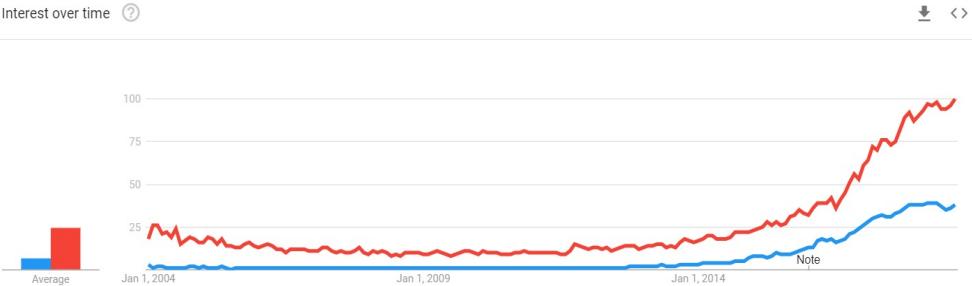
Deep Learning's Promise To The Marketing Sector

REPORT
All-optical machine learning using diffractive deep neural networks

MarketWatch

PRESS RELEASE
MEDIA ALERT: Deep Learning Acceleration Startup Mipsology to Exhibit at Xilinx Developer Forum

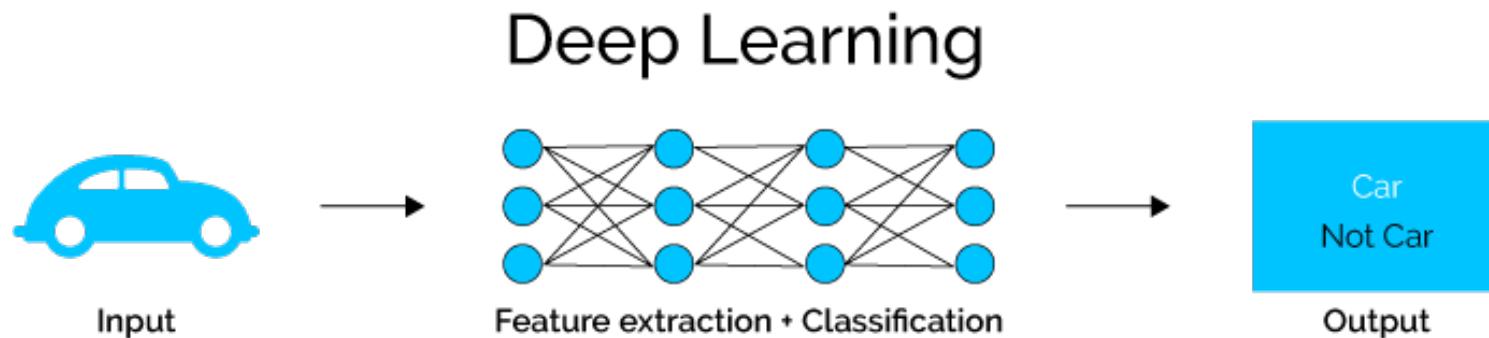
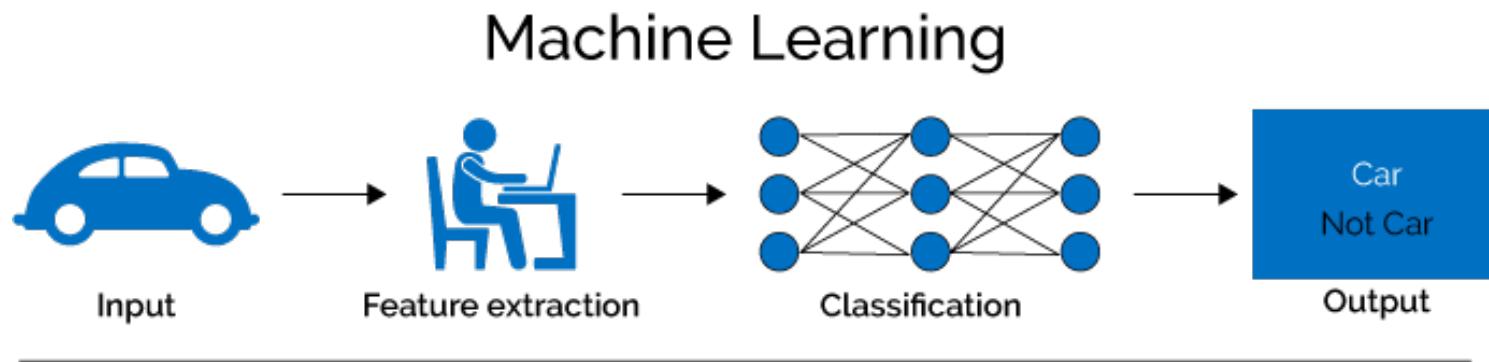
Deep learning for classifying fibrotic lung disease on high-resolution computed tomography: a case-cohort study

Interest over time 

- machine learning
- deep learning

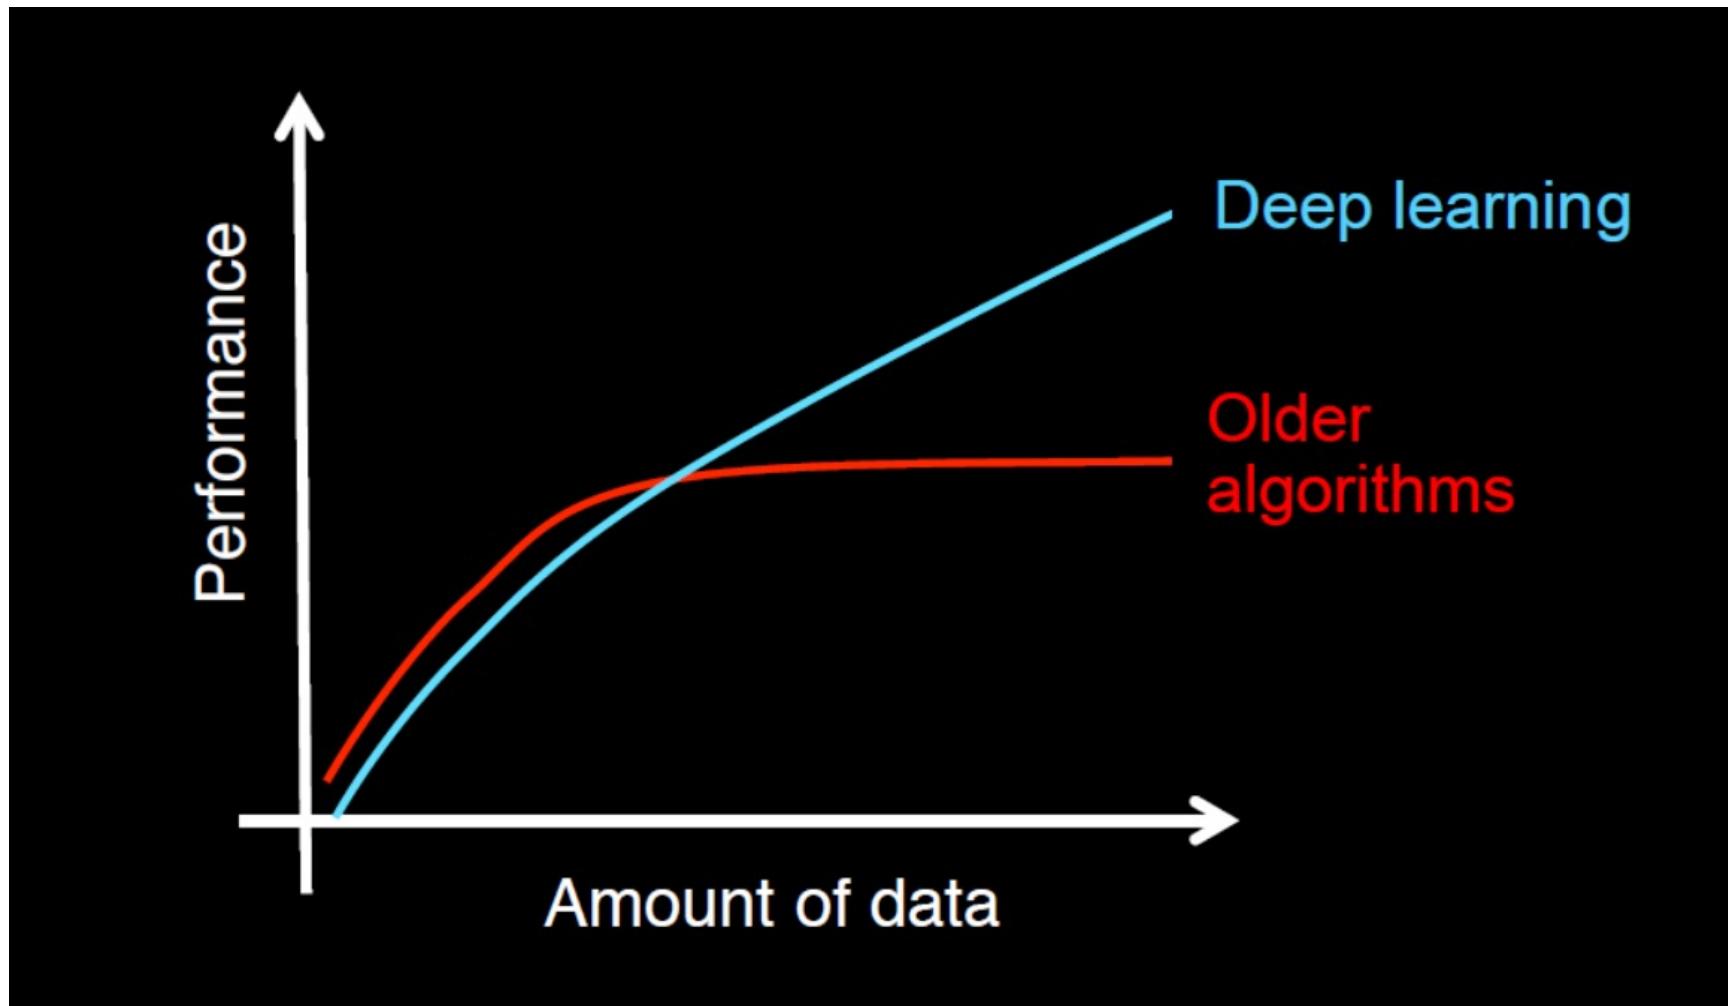
Introduction

- ◆ Deep learning: **end-to-end multi-layer** neural network for feature generation and processing



[<https://www.xenonstack.com/blog/static/public/uploads/media/machine-learning-vs-deep-learning.png>]

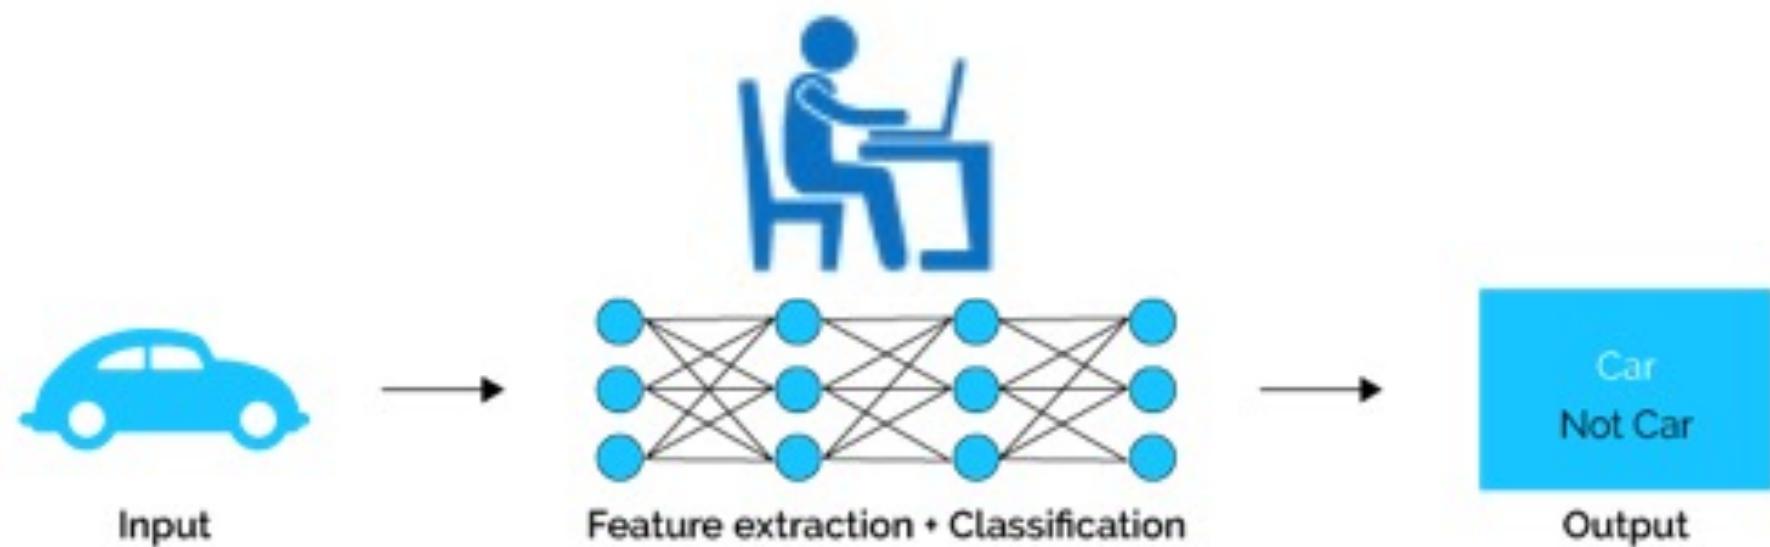
Introduction



[Andrew Ng, <http://cs229.stanford.edu/materials/CS229-DeepLearning.pdf>]

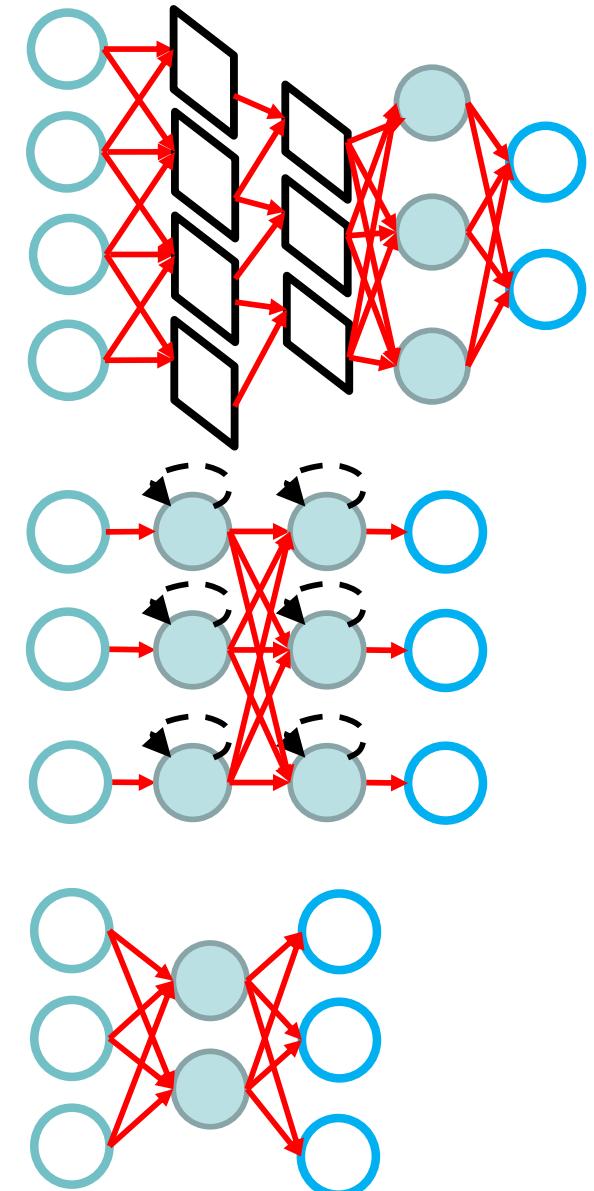
Introduction

Need to design architectures!



Introduction

- ◆ Deep Learning Architecture
- ◆ Supervised
 - Convolutional neural network (CNN)
 - Recurrent neural network (RNN)
 - Multilayer perceptron
- ◆ Unsupervised
 - Autoencoder
 - Boltzman machine
- ◆ Generative Adversarial Networks



Fundamentals of Deep Learning

What's Deep Learning?

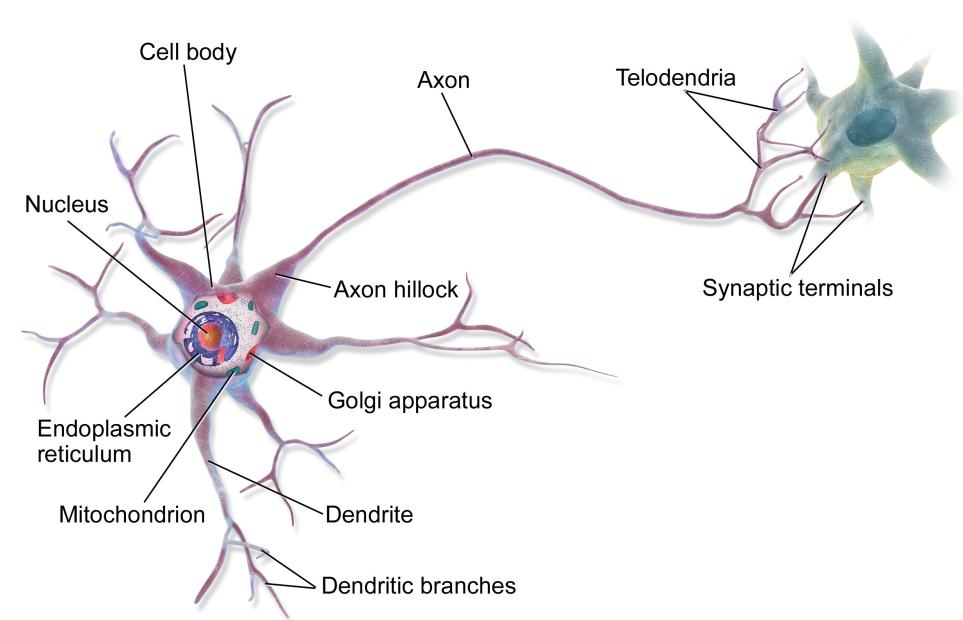
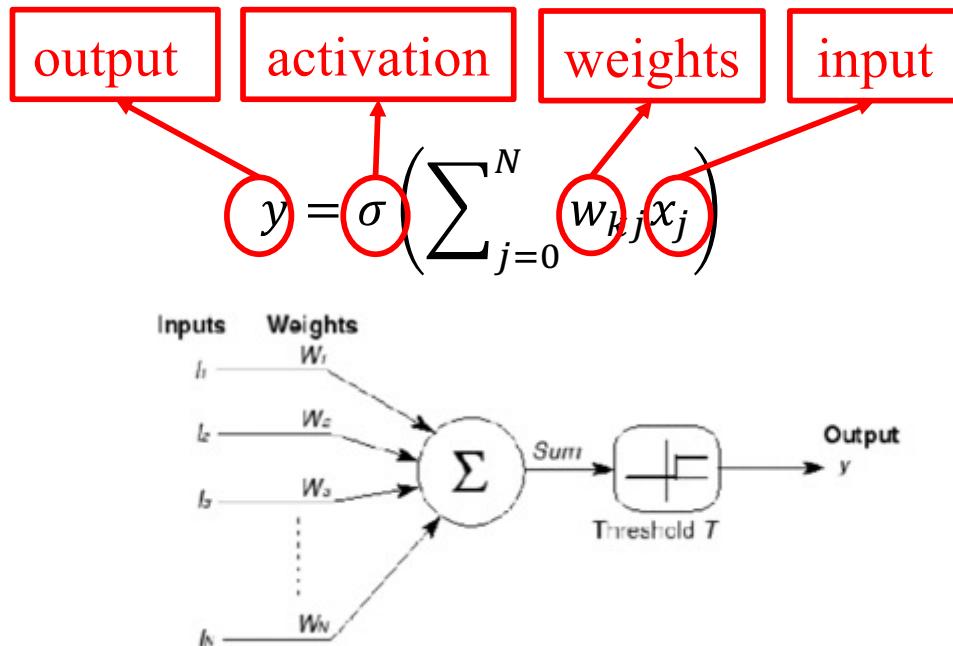
- ◆ Deep learning is part of a broader family of machine learning methods based on the layers used in **artificial neural networks**.
- ◆ Artificial neural networks (ANNs) are computing systems that are inspired by, but not identical to, biological neural networks that constitute animal brains.



Neuron

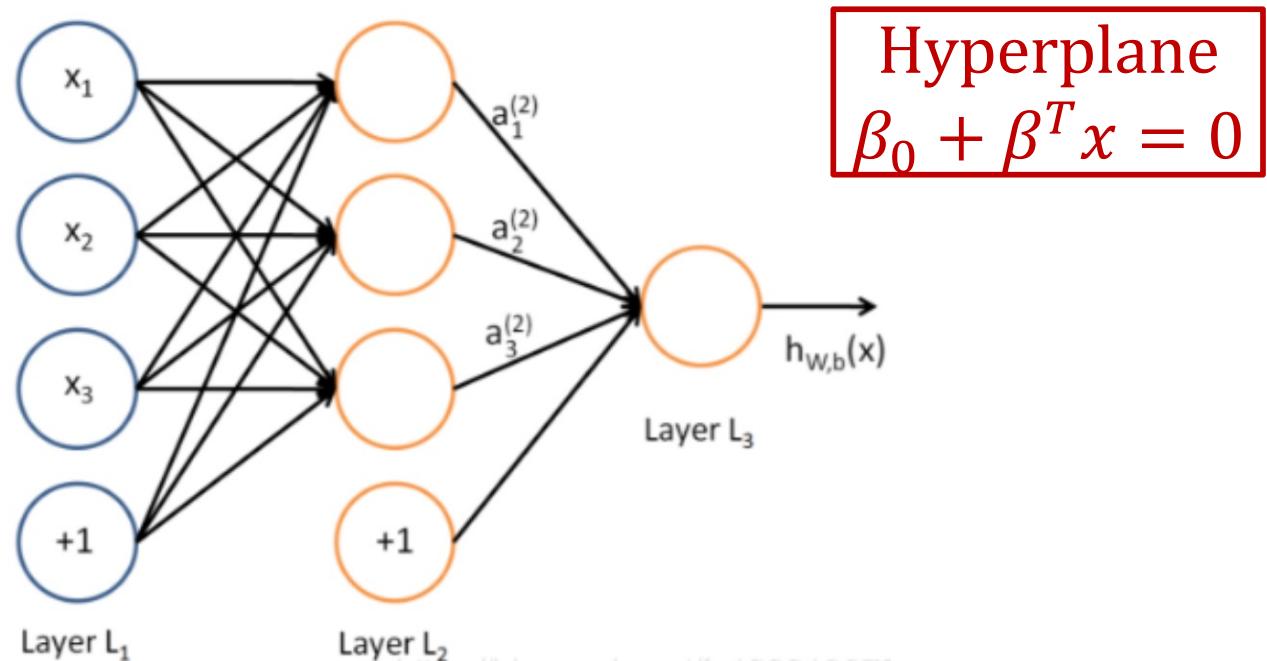
◆ The MCP (McCulloch-Pitts) neuron

- First proposed in 1943 by the neurophysiologist Walter S. McCulloch and the logician Walter Pitts, the MCP neuron is a simple mathematical model of a *biological neuron*.



Perceptron

- ◆ The perceptron algorithm was invented in 1958 by Rosenblatt used to binary classification.
- ◆ Perceptron trained with gradient descent algorithm, and its convergence was guaranteed in 1962.

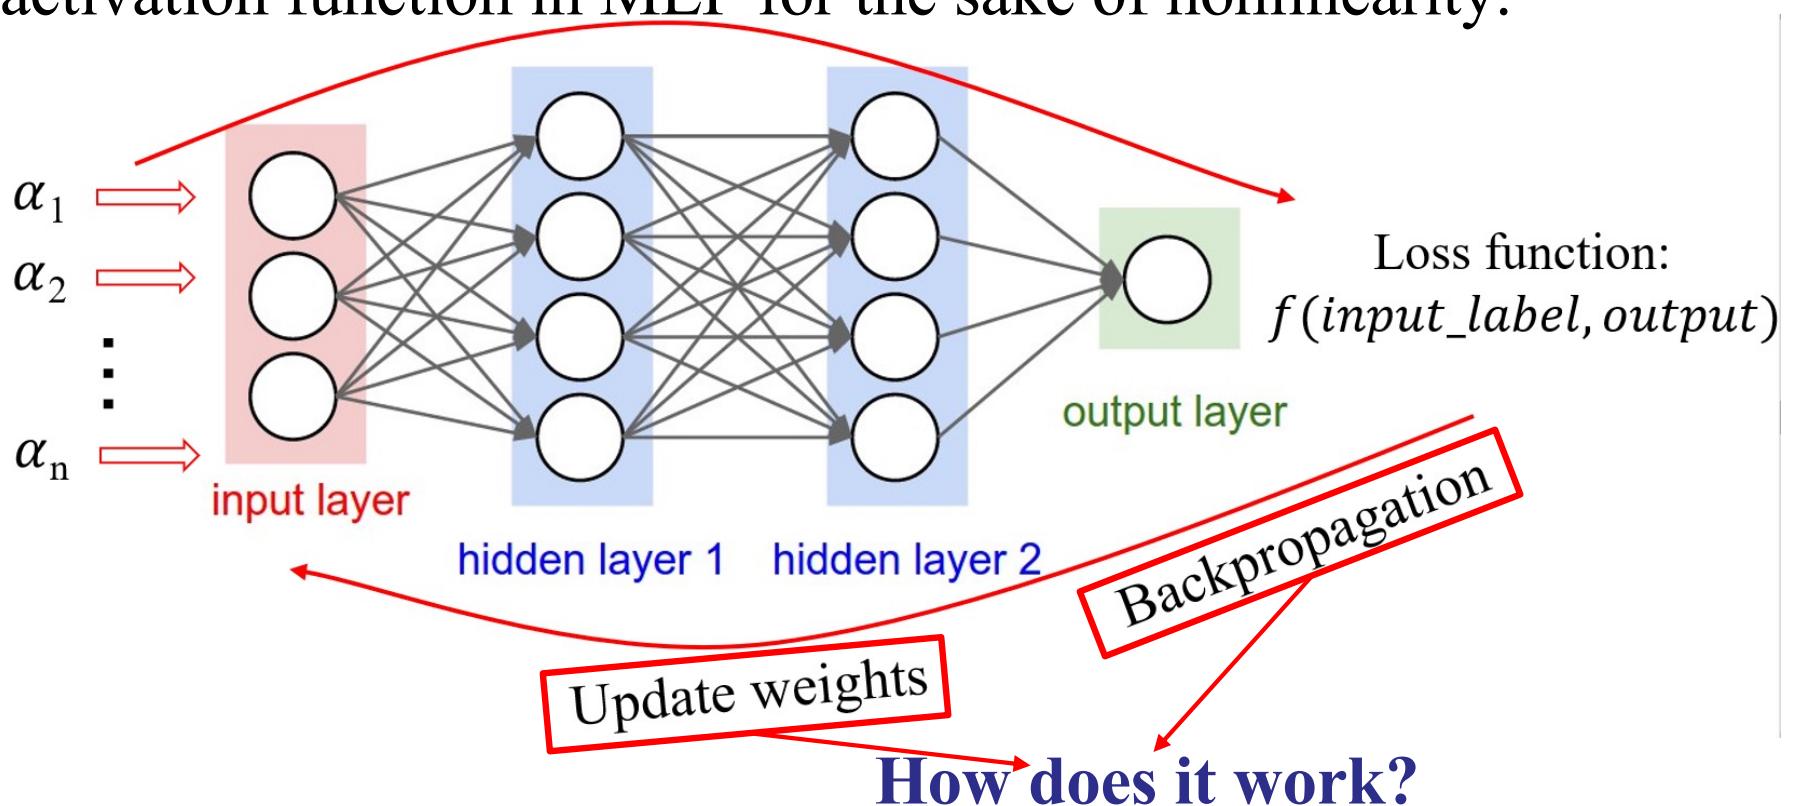


Limitation of Perceptron

- ◆ Marvin Minsky proved that perceptron is a linear model essentially in 1969, and perceptron cannot handle nonlinear classification.
- ◆ To solve this problem, it is necessary to increase the number of layers.
- ◆ Training algorithm for a multi-layer perceptron (MLP)

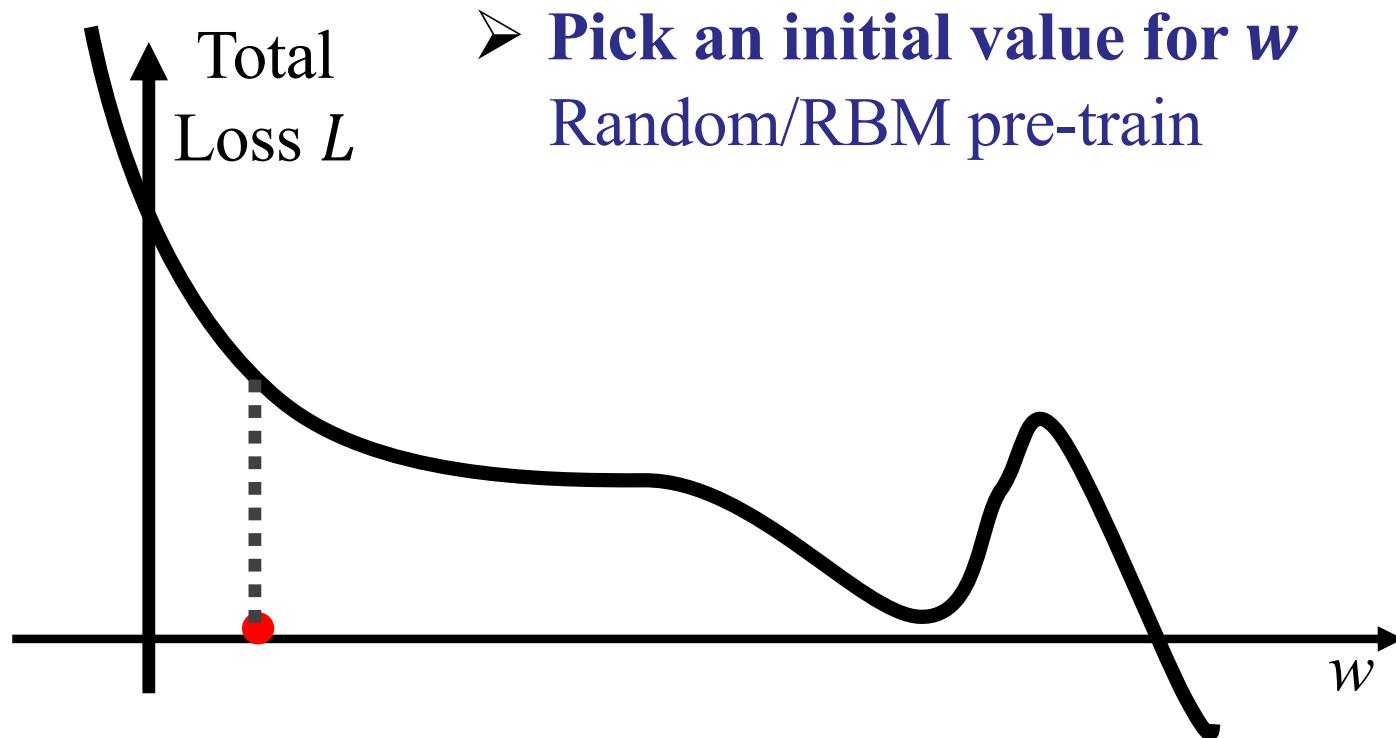
Backpropagation

- ◆ Geoffrey Hinton proposed the Backpropagation (BP) algorithm in 1986 for training the MLP.
- ◆ The sigmoid function $S(x) = 1/(1 + e^{-x})$ is used as the activation function in MLP for the sake of nonlinearity.



Gradient Descent

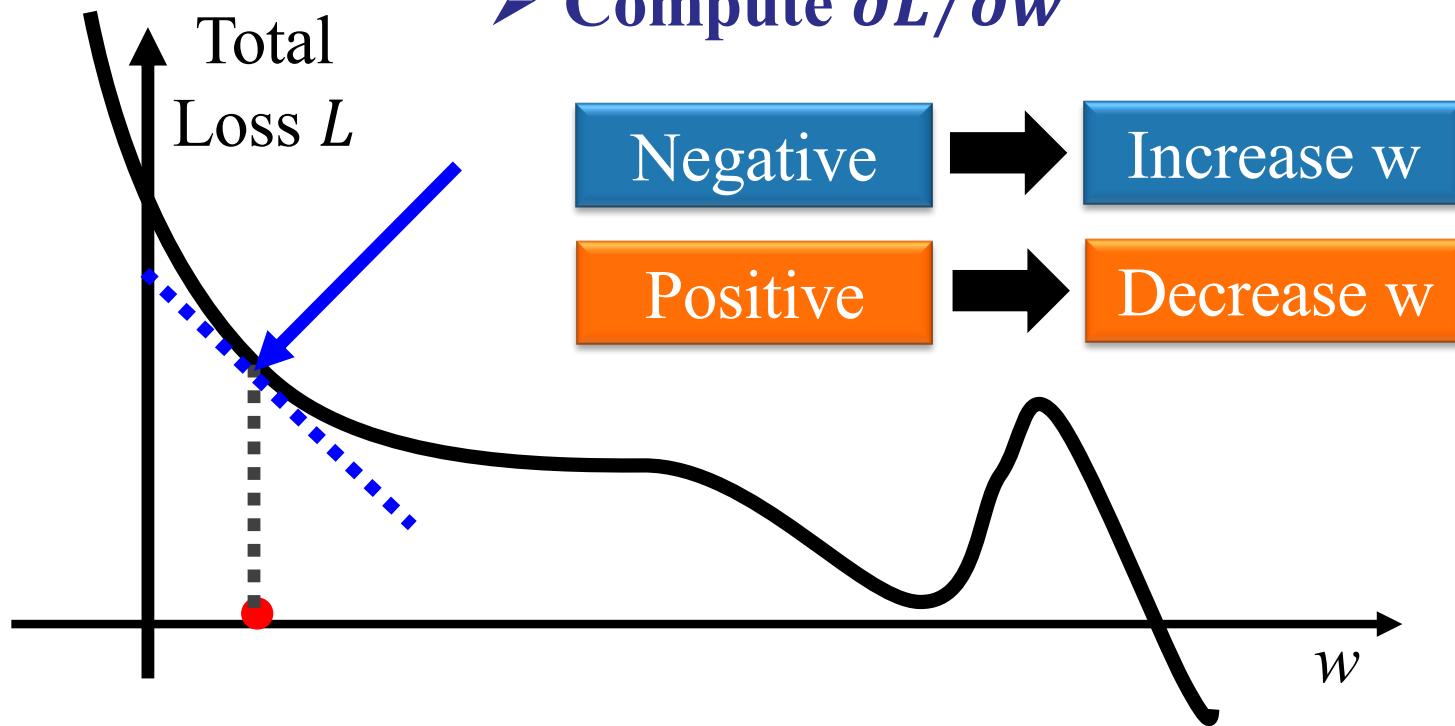
Network parameters $\theta = \{w_1, w_2, w_3, \dots, b_1, b_2, b_3, \dots\}$



Gradient Descent

Network parameters $\theta = \{w_1, w_2, w_3, \dots, b_1, b_2, b_3, \dots\}$

- Pick an initial value for w
- Compute $\partial L / \partial w$



Gradient Descent

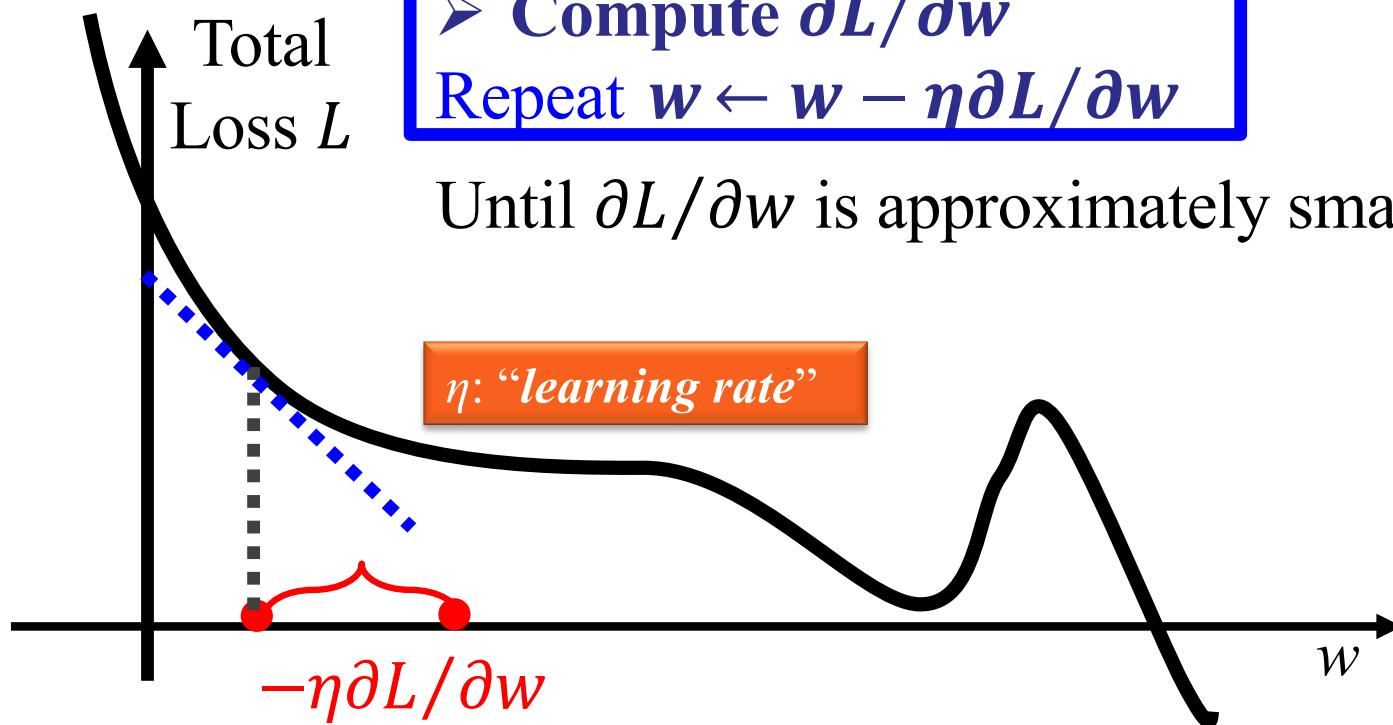
Network parameters $\theta = \{w_1, w_2, w_3, \dots, b_1, b_2, b_3, \dots\}$

➤ Pick an initial value for w

➤ Compute $\partial L / \partial w$

Repeat $w \leftarrow w - \eta \partial L / \partial w$

Until $\partial L / \partial w$ is approximately small

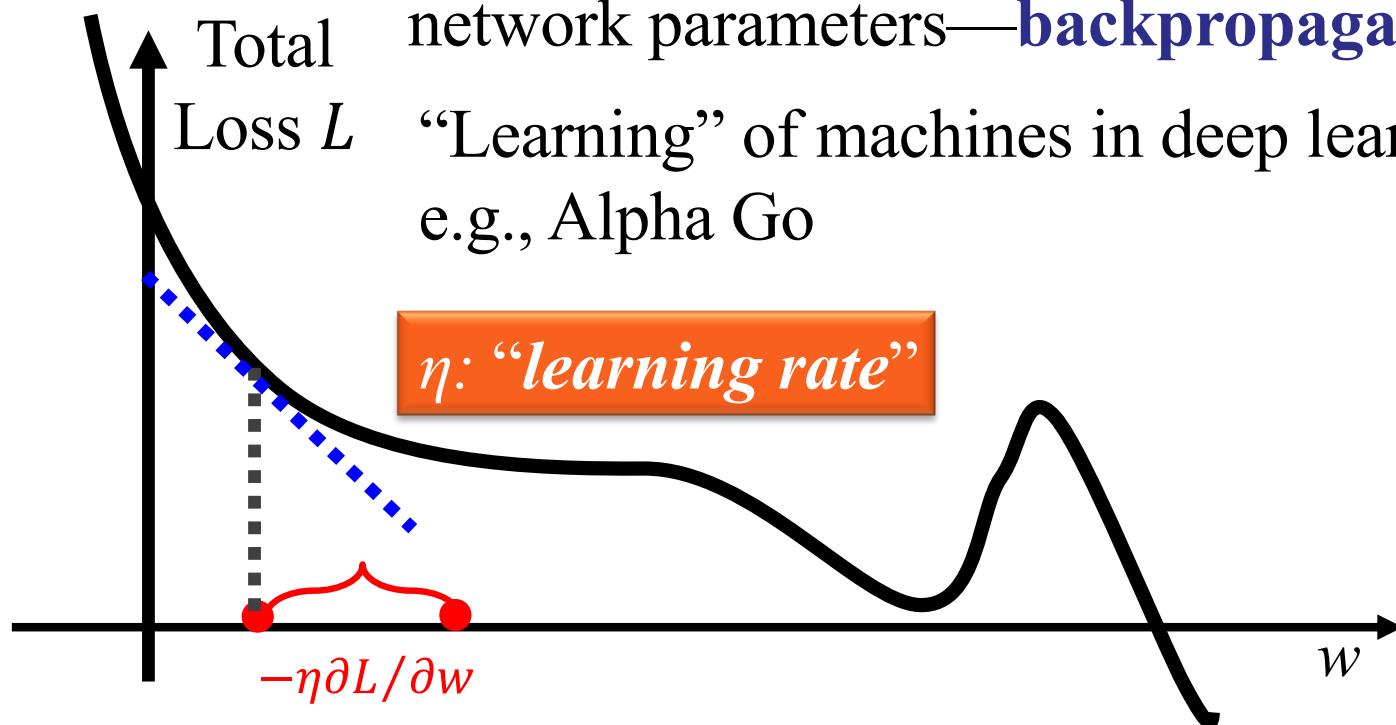


Gradient Descent

Network parameters $\theta = \{w_1, w_2, w_3, \dots, b_1, b_2, b_3, \dots\}$

An efficient way to compute the $\partial L / \partial w$ of the network parameters—**backpropagation**.

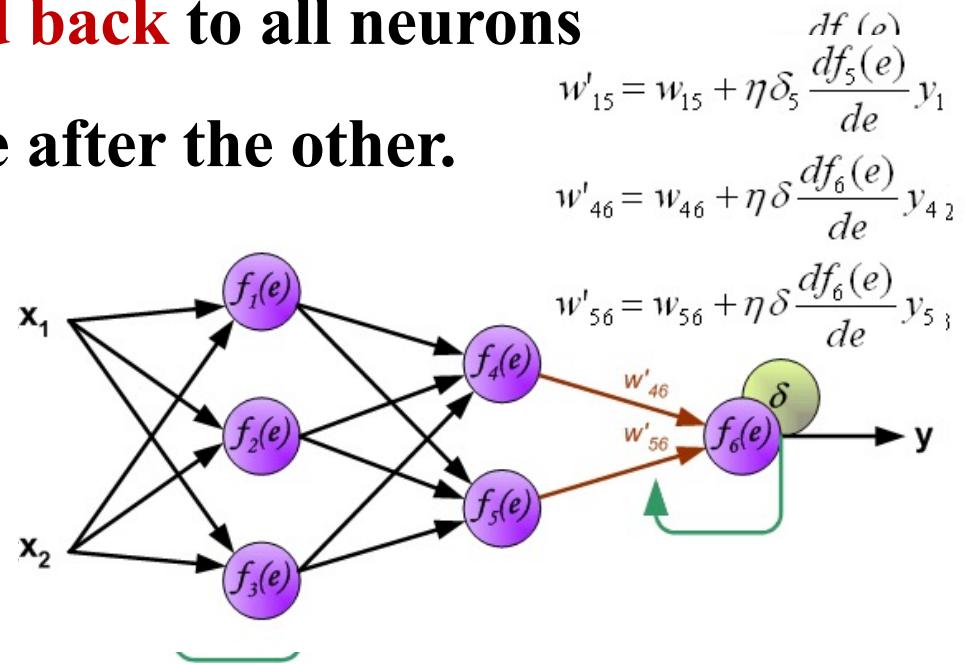
Total Loss L “Learning” of machines in deep learning,
e.g., Alpha Go



Backpropagation

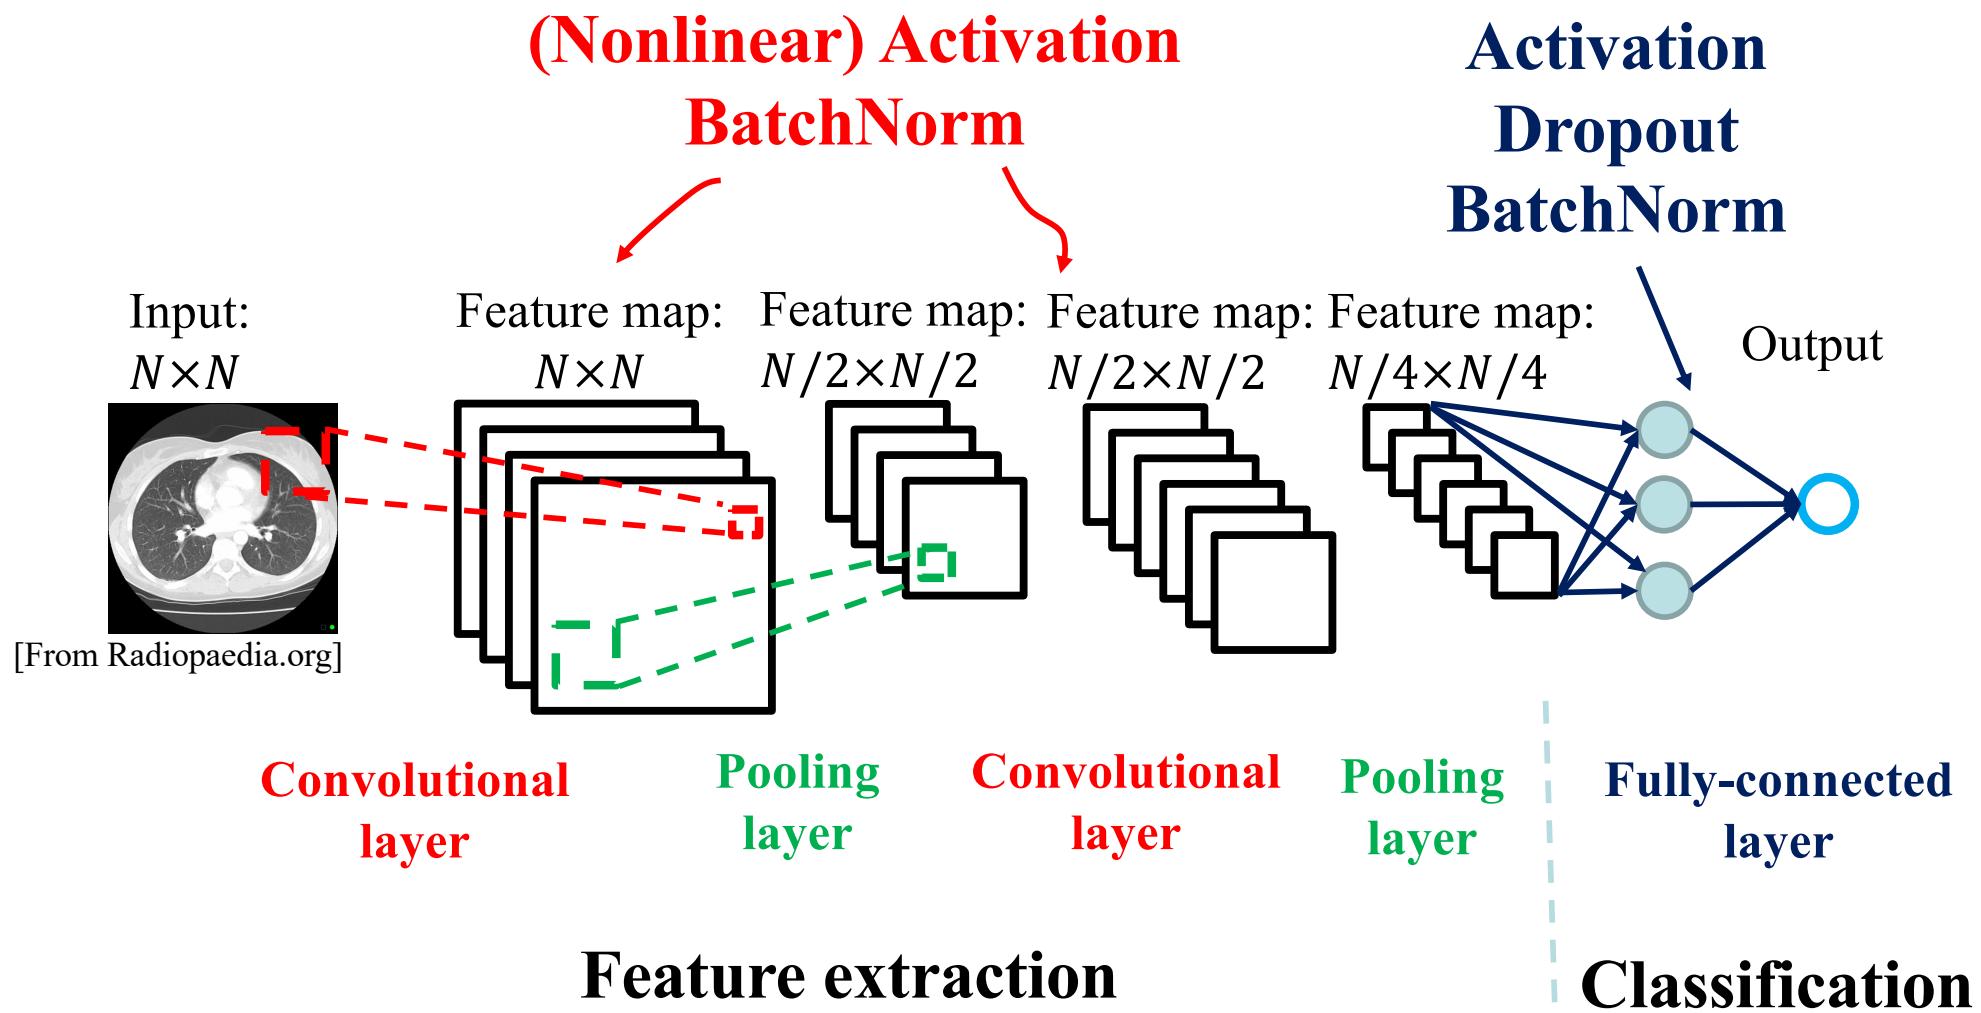
- ◆ Three-layer neural network with two inputs and one output

- Forward propagation through the hidden & output layers
- Compare the output y with the target z (label).
- The error δ is propagated back to all neurons from output to inputs one after the other.



Convolutional Neural Networks (CNNs)

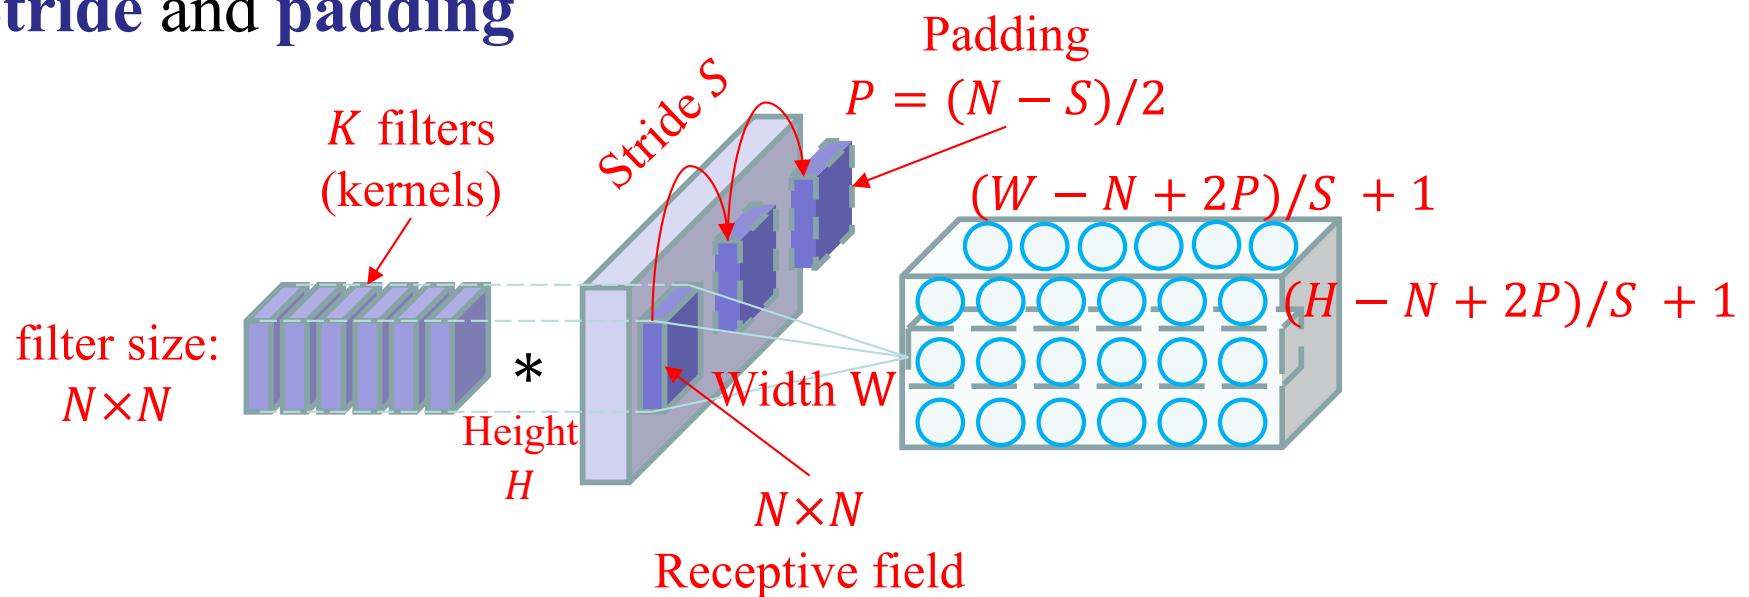
Convolutional Neural Networks



Convolutional Neural Networks

◆ Convolutional layer

- A set of **filters (kernels)** learned to detect specific types of features at specific spatial location
- **Receptive field** and **neuron**
- **Stride** and **padding**

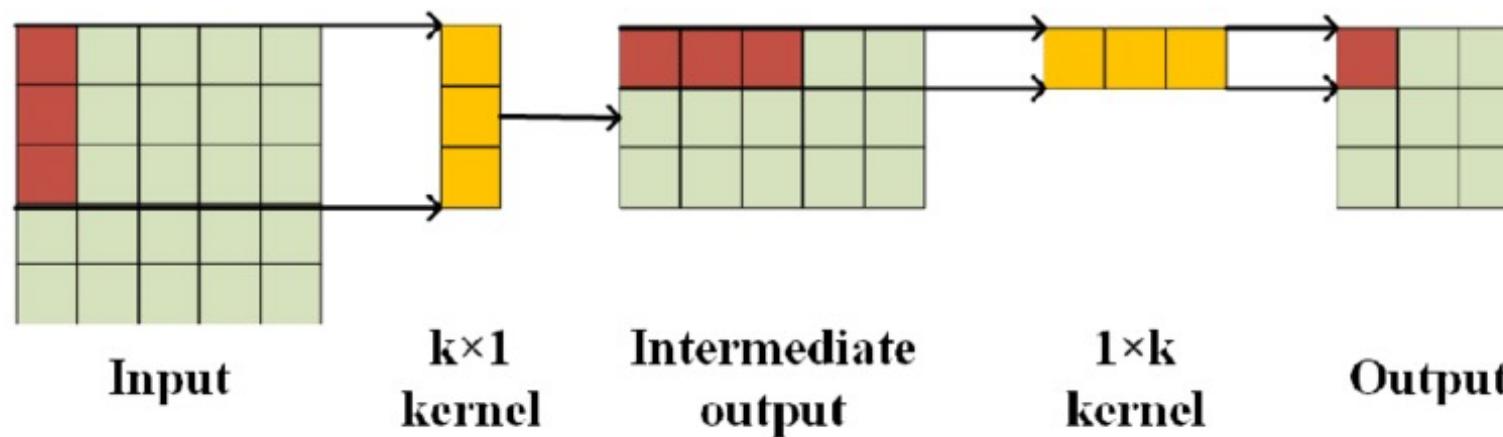


Convolutional Neural Networks

◆ Separable Convolutions

◆ Spatial separable convolutions

Transform convolutions with $K \times K$ kernels into convolutions with $K \times 1$ and $1 \times K$ kernels



Convolutional Neural Networks

◆ Separable Convolutions

◆ Depthwise separable convolutions

Realize convolutions with $N \times N$ kernels for CI input channels and CO output channels with depthwise convolutions with CI $N \times N \times 1$ kernels and pointwise convolution with CO $1 \times 1 \times CI$ kernels

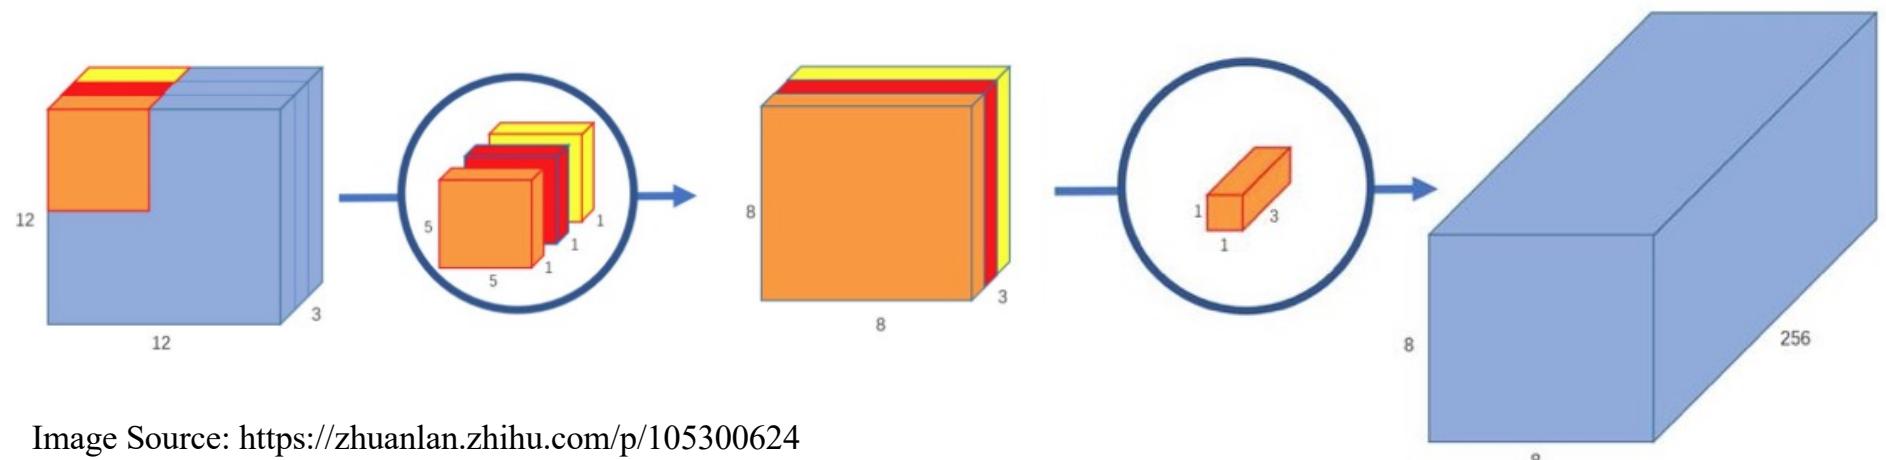
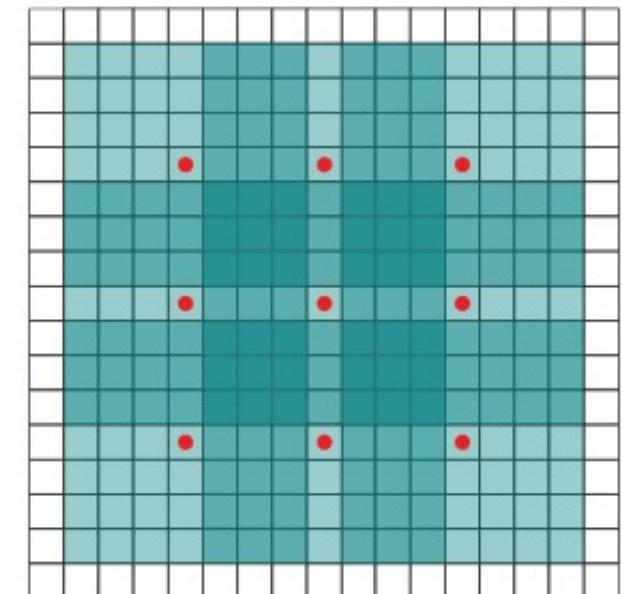
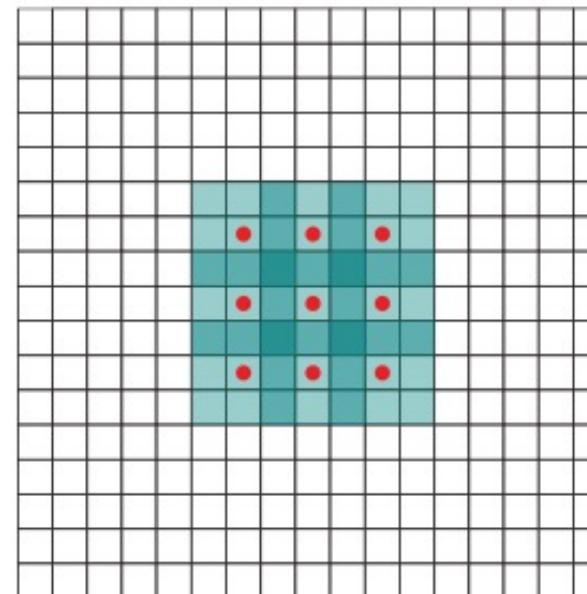
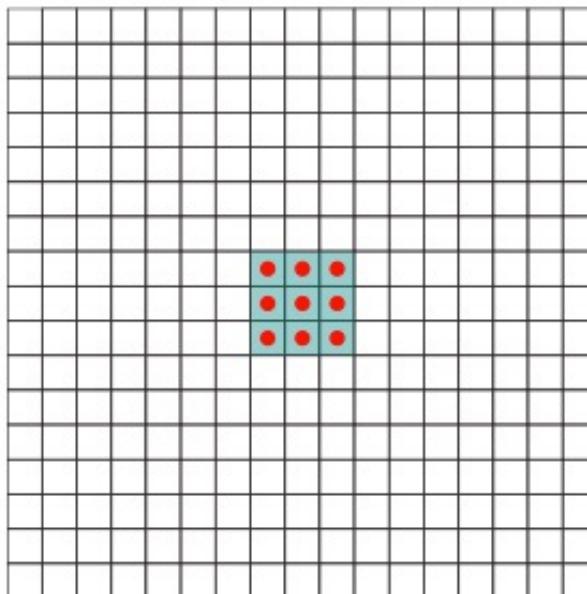


Image Source: <https://zhuanlan.zhihu.com/p/105300624>

Convolutional Neural Networks

◆ Dilated Convolutions

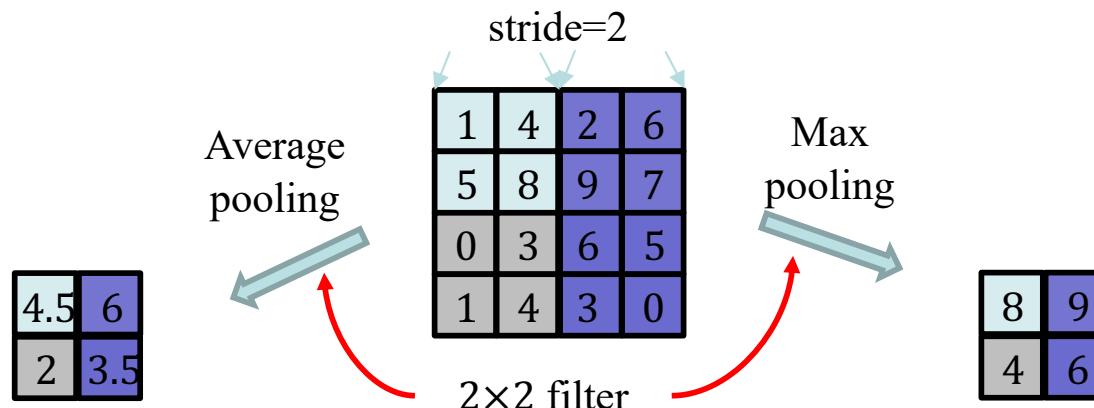
- ◆ Exponentially expand receptive fields without losing resolution or coverage
- ◆ Maintain the size of feature maps without excessive computations



Convolutional Neural Networks

◆ Pooling layer

- **Nonlinear downsampling** to reduce spatial size and number of parameters
- Suppress overfitting
- Common routines: **max pooling** & **average pooling**



Convolutional Neural Network

◆ Activation layer

- Introduce nonlinearities to linear convolutional layers
- Nonlinear functions: tanh, sigmoid, ReLU (rectified linear units)

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\text{relu}(x) = \max(x, 0)$$

- ReLU allows faster training without significant difference in accuracy

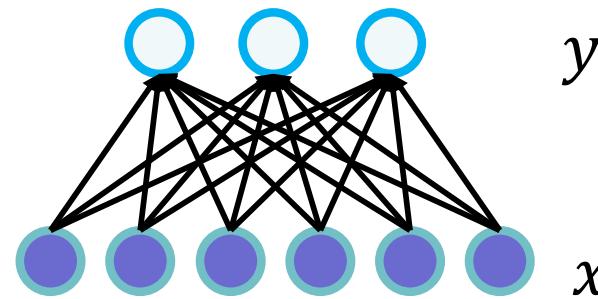
Convolutional Neural Network

◆ Fully (densely) connected layer

- Each unit is connected to all units of the previous layers

$$y = \sigma(wx + b)$$

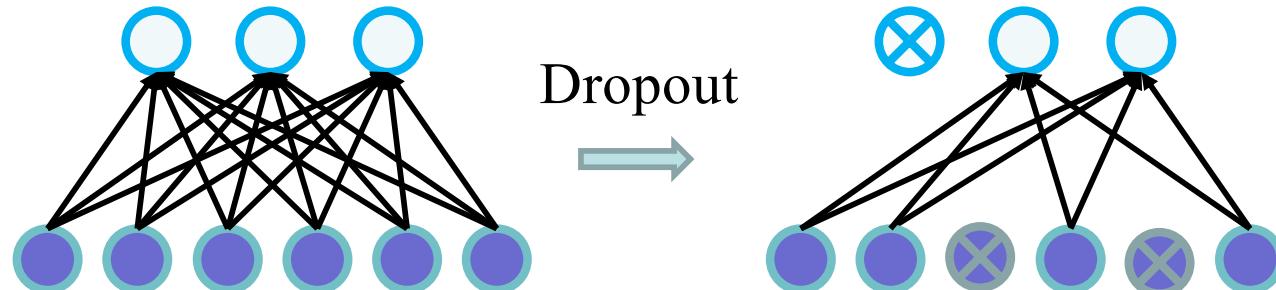
- Learn nonlinear combinations of features for classification
- Activation function to introduce nonlinearity



Convolutional Neural Network

◆ Dropout: Regularization technique to reduce overfitting

- An efficient way to approximately combine different neural network architectures
- Randomly drop a portion of nodes and their connections with a probability (dropout rate) for gradient computation



Convolutional Neural Network

◆ Batch Normalization

- ◆ Allow higher learning rates and robust to initialization
(suppress explosive or vanishing gradients and early stop in local minimum)
- ◆ Regularize the deep learning model
- ◆ Normalize by scaling and shifting activations
- ◆ For minibatch $B = \{x_1, \dots, x_m\}$,

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i, \quad \sigma_B^2 = \frac{1}{m} \sum_i (x_i - \mu_B)^2$$
$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}, \quad y_i = \gamma \hat{x}_i + \beta$$

Convolutional Neural Network

- ## ◆ Layer-by-layer training with backpropagation

- For convolutional layer l , forward process:

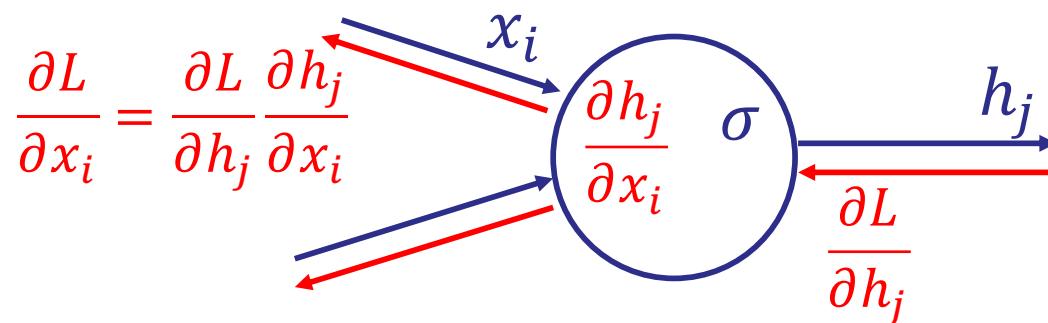
$$x_{i,j}^l = \sum_{u=0}^{n-1} \sum_{v=0}^{n-1} w_{u,v}^l h_{i+u, j+v}^{l-1} + b_{i,j}^l, \quad h_{i,j}^l = f(x_{i,j}^l)$$

weights

activation function

- ## ■ Backward process:

$$\frac{\partial L}{\partial w_{i,j}^l} = \sum_{u=0}^{H-n} \sum_{v=0}^{W-n} \delta_{u,v}^l n_{i+u, j+v}^{l-1}$$

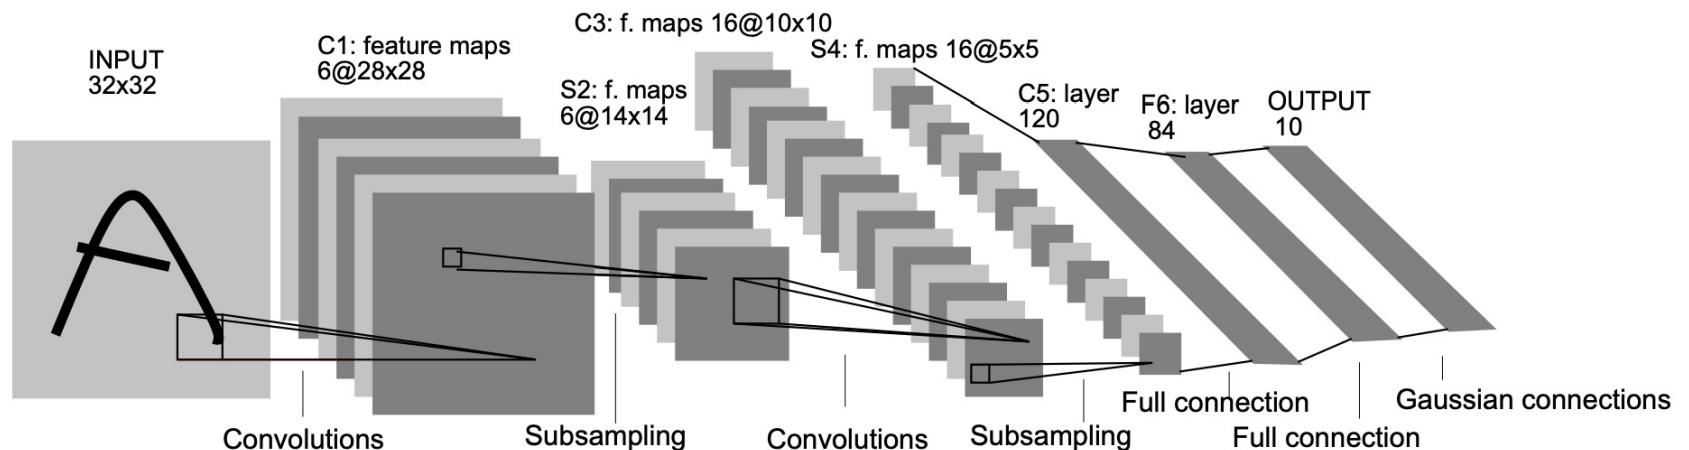


Examples of CNN

◆ LeNet: The pioneer of CNN

◆ Insights:

- Local Receptive Fields
- Shared Weights
- Sub-Sampling



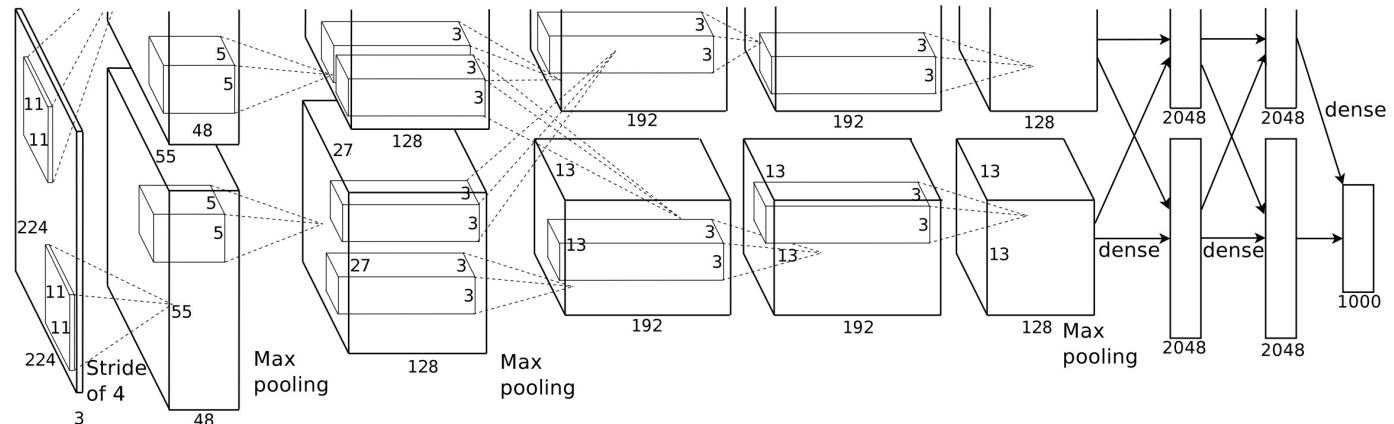
LeCun, Yann, et al. "Gradient-based learning applied to document recognition." *Proceedings of the IEEE* 86.11 (1998): 2278-2324.

Examples of CNN

◆ AlexNet: The Revealer of New Era

◆ Insights:

- ReLU
- Dropout
- Overlapping Pooling
- Groupwise Convolution
- Local Response Normalization



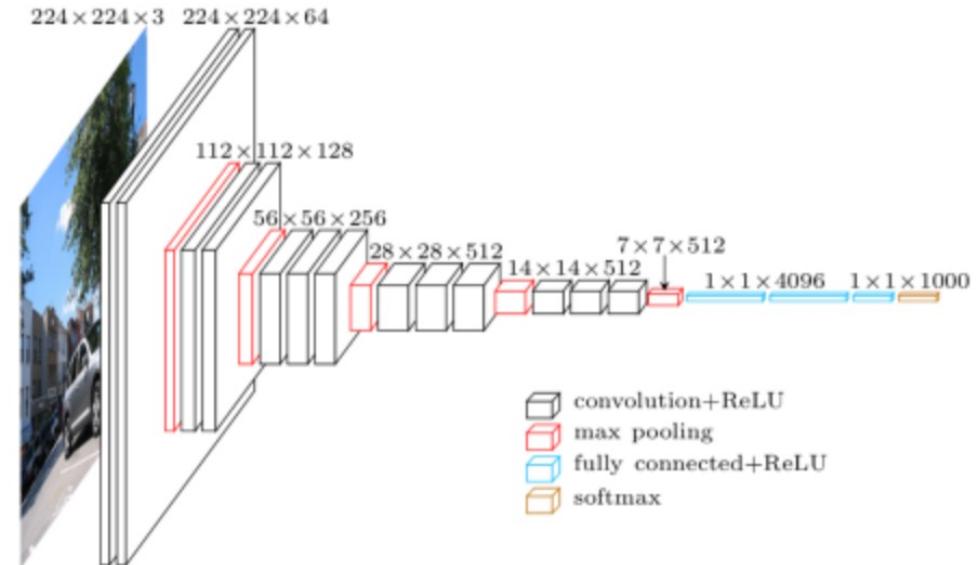
Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.

Examples of CNN

- ◆ VGG: Very Deep CNN

- ◆ Insights:

- Substitute large convolution kernels with Multi-layer small convolution kernels
- Replace FC with convolution layer in evaluation
- Expand channels
- Deeper and Wider



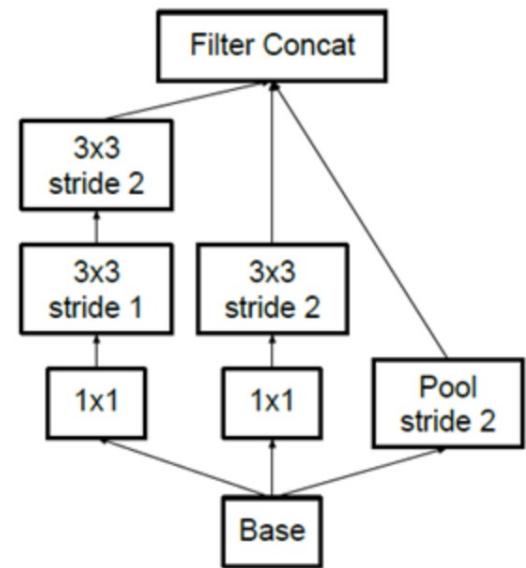
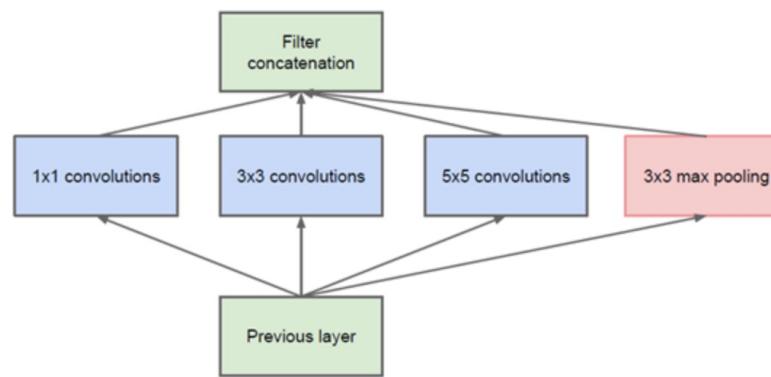
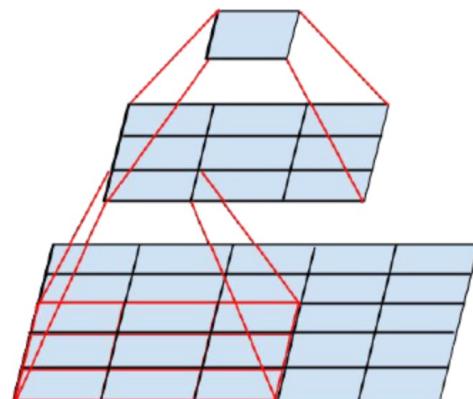
Simonyan K, Zisserman A. Very deep convolutional networks for large-scale image recognition[J]. arXiv preprint arXiv:1409.1556, 2014.

Examples of CNN

◆ InceptionNet: The Champion of ILSVRC14

◆ Insights:

- Multi-scale kernel size
- Factorizing Convolutions
- Downsampling in parallel



Szegedy C, Liu W, Jia Y, et al. Going deeper with convolutions[C]//Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition. 2015: 1-9.

Examples of CNN

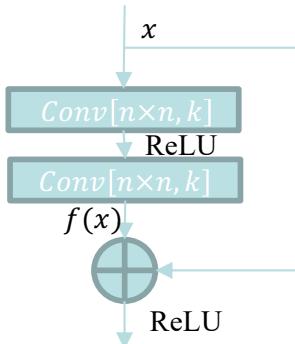
◆ ResNet

◆ Insights:

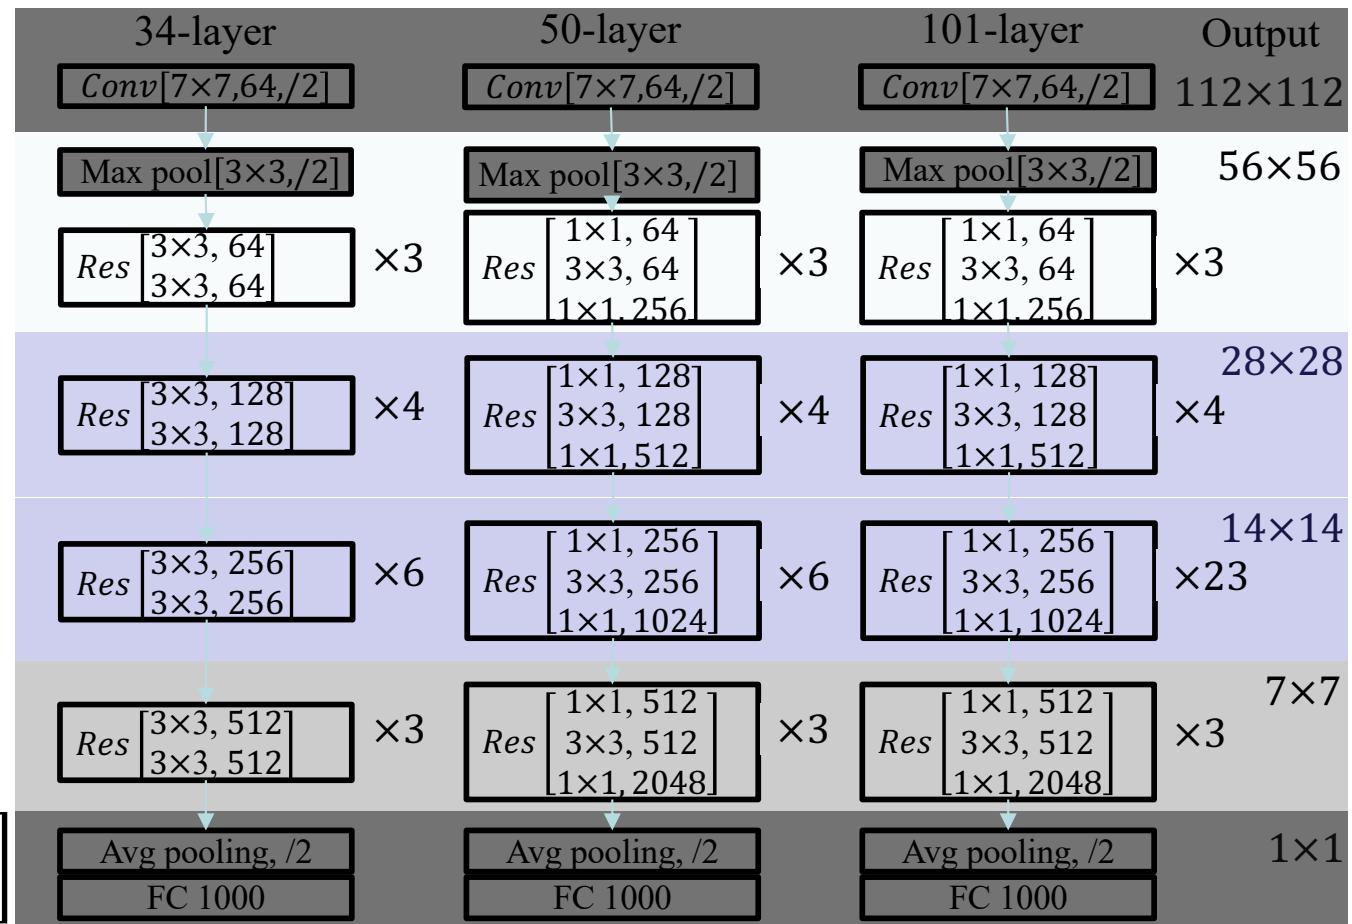
- Train residual functions
- Shortcut connection

$$y = f(x) + x$$

$$f = w_2 \sigma(w_1 x)$$



$$\text{Res} \begin{bmatrix} n \times n, k \\ n \times n, k \end{bmatrix}$$



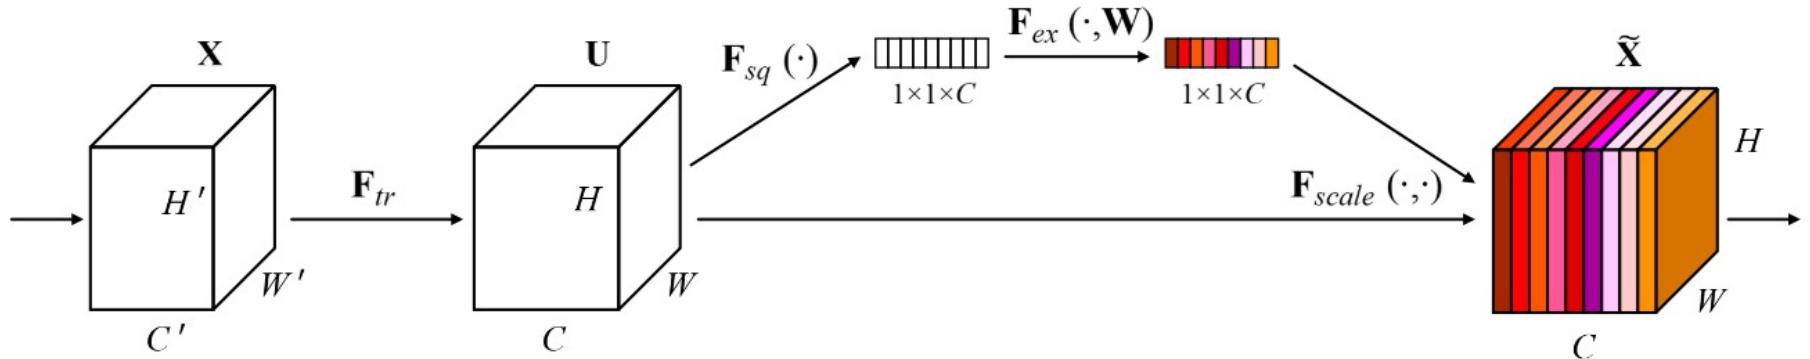
K. He, X. Zhang, S. Ren, J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 770-778.

Examples of CNN

- ◆ **SENet: The Champion of ILSVRC17**

- ◆ **Insights:**

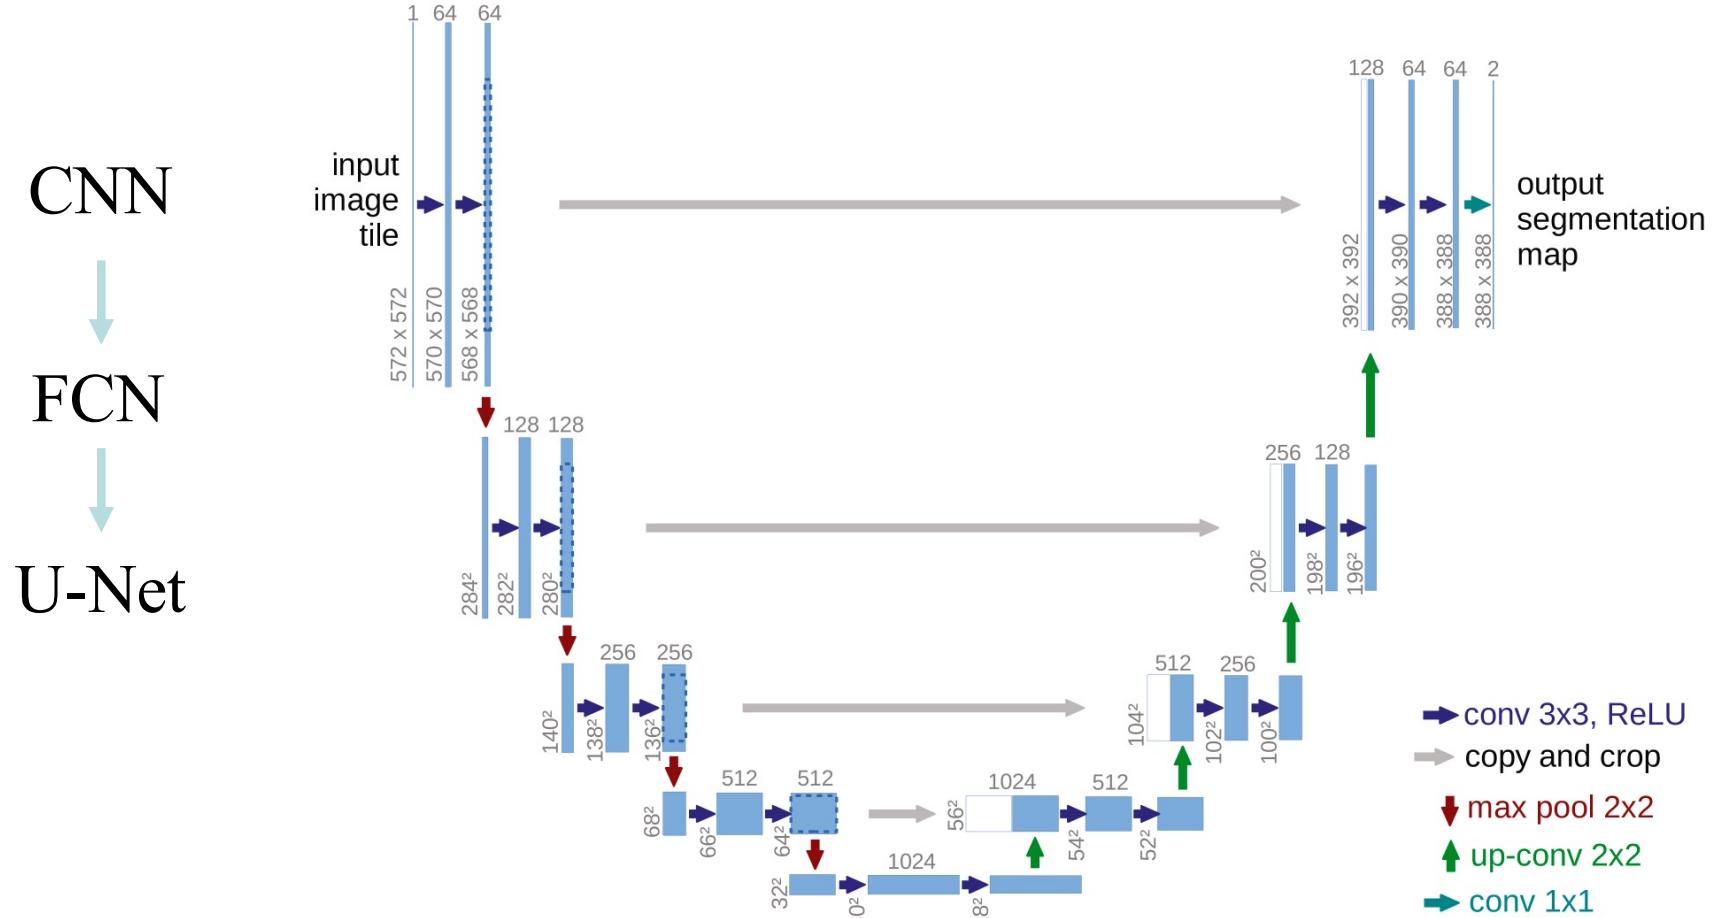
- Applying channel attention mechanism to CNN.
- Squeeze & Excitation



Hu J, Shen L, Sun G. Squeeze-and-excitation networks[C]//Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition. 2018: 7132-7141.

Examples of CNN

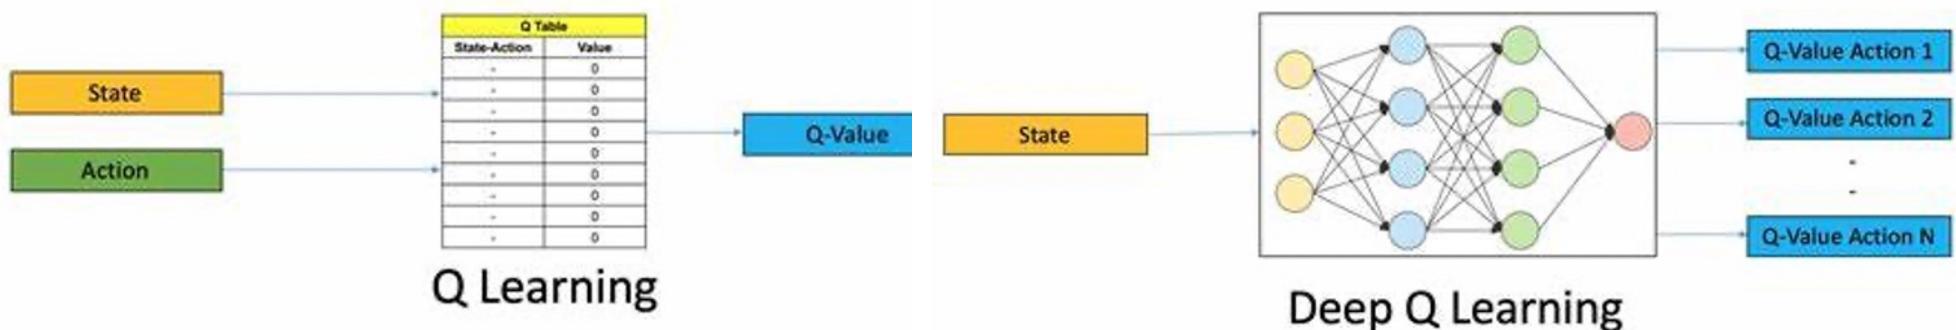
◆ U-Net: The milestone in medical image segmentation



Example: Deep Q-Network

- ◆ Neural network to approximate the Q-value function.
 - ◆ Input: the state
 - ◆ Output: the Q-value of all possible actions
- ◆ Maximum output of the Q-network as the next action
- ◆ Bellman equation for Q-value update

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$



Recurrent Neural Networks (RNNs)

Recurrent Neural Networks

- ◆ Recurrent networks introduce cycles and a notion of time

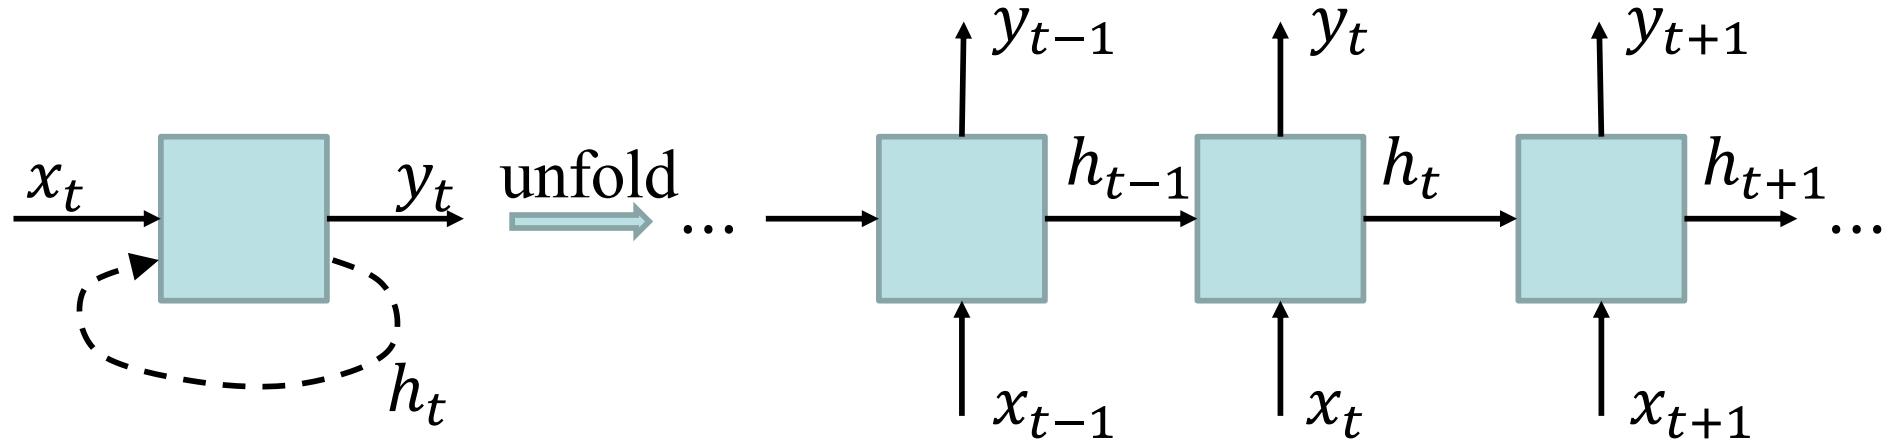
$$h_t = \sigma(x_t w_{xh} + h_{t-1} w_{hh} + b_h)$$

$$y_t = \sigma(h_t w_{hy} + b_y)$$

- ◆ Map a sequence of inputs x_1, \dots, x_N to sequence of outputs

$$y_1, \dots, y_n$$

- ◆ Direct acyclic graph (DAG) for back propagation

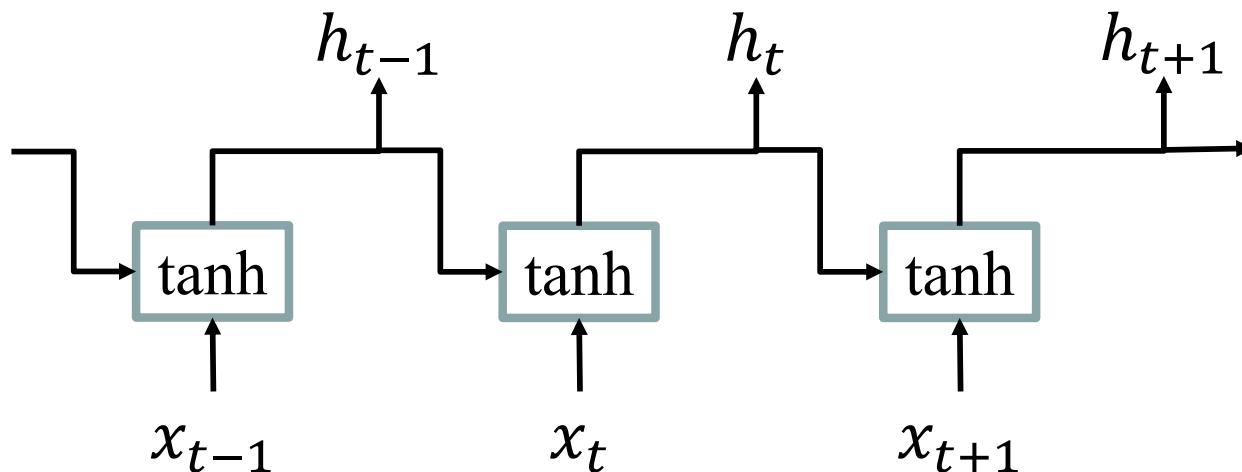


Recurrent Neural Networks

◆ Standard RNNs

$$h_t = \tanh(x_t w_{xh} + h_{t-1} w_{hh} + b_h)$$

$$\hat{y}_t = \text{softmax}(h_t w_{hy} + b_y)$$



Recurrent Neural Networks

$$\diamond L(y, \hat{y}) = - \sum_t L_t(y_t, \hat{y}_t)$$

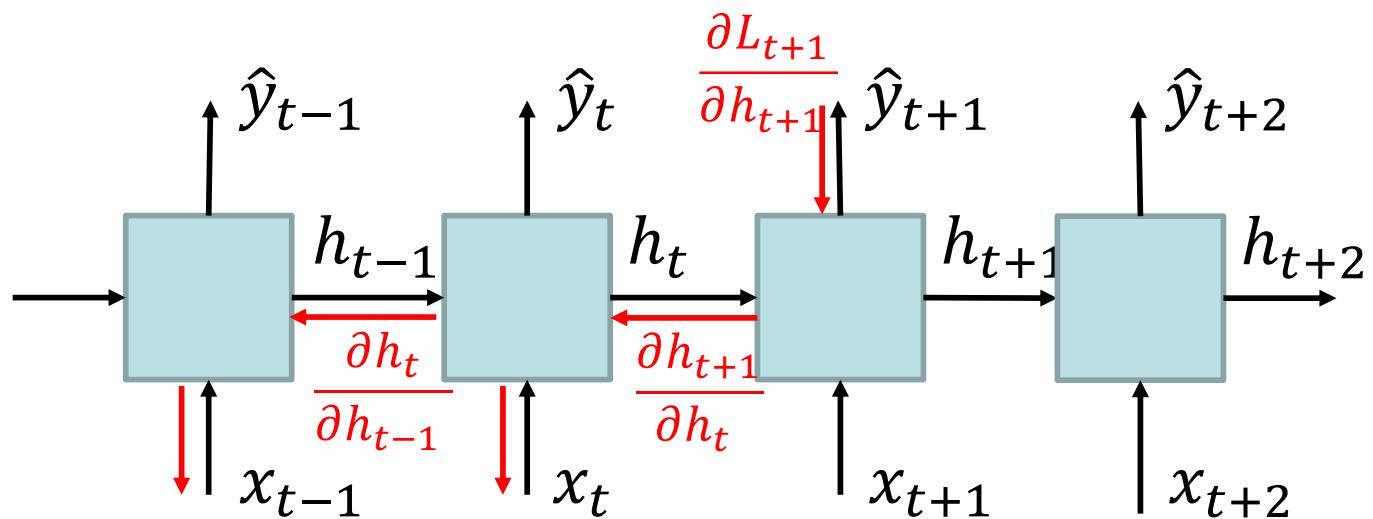
Label

$$\diamond \frac{\partial L_t}{\partial w} = \frac{\partial L_t}{\partial \hat{y}_t} \frac{\partial \hat{y}_t}{\partial h_t} \frac{\partial h_t}{\partial w} = \sum_{i=1}^t \frac{\partial L_t}{\partial \hat{y}_t} \frac{\partial \hat{y}_t}{\partial h_t} \frac{\partial h_t}{\partial h_i} \frac{\partial h_i}{\partial w}$$

Depends on h_{t-1} and w

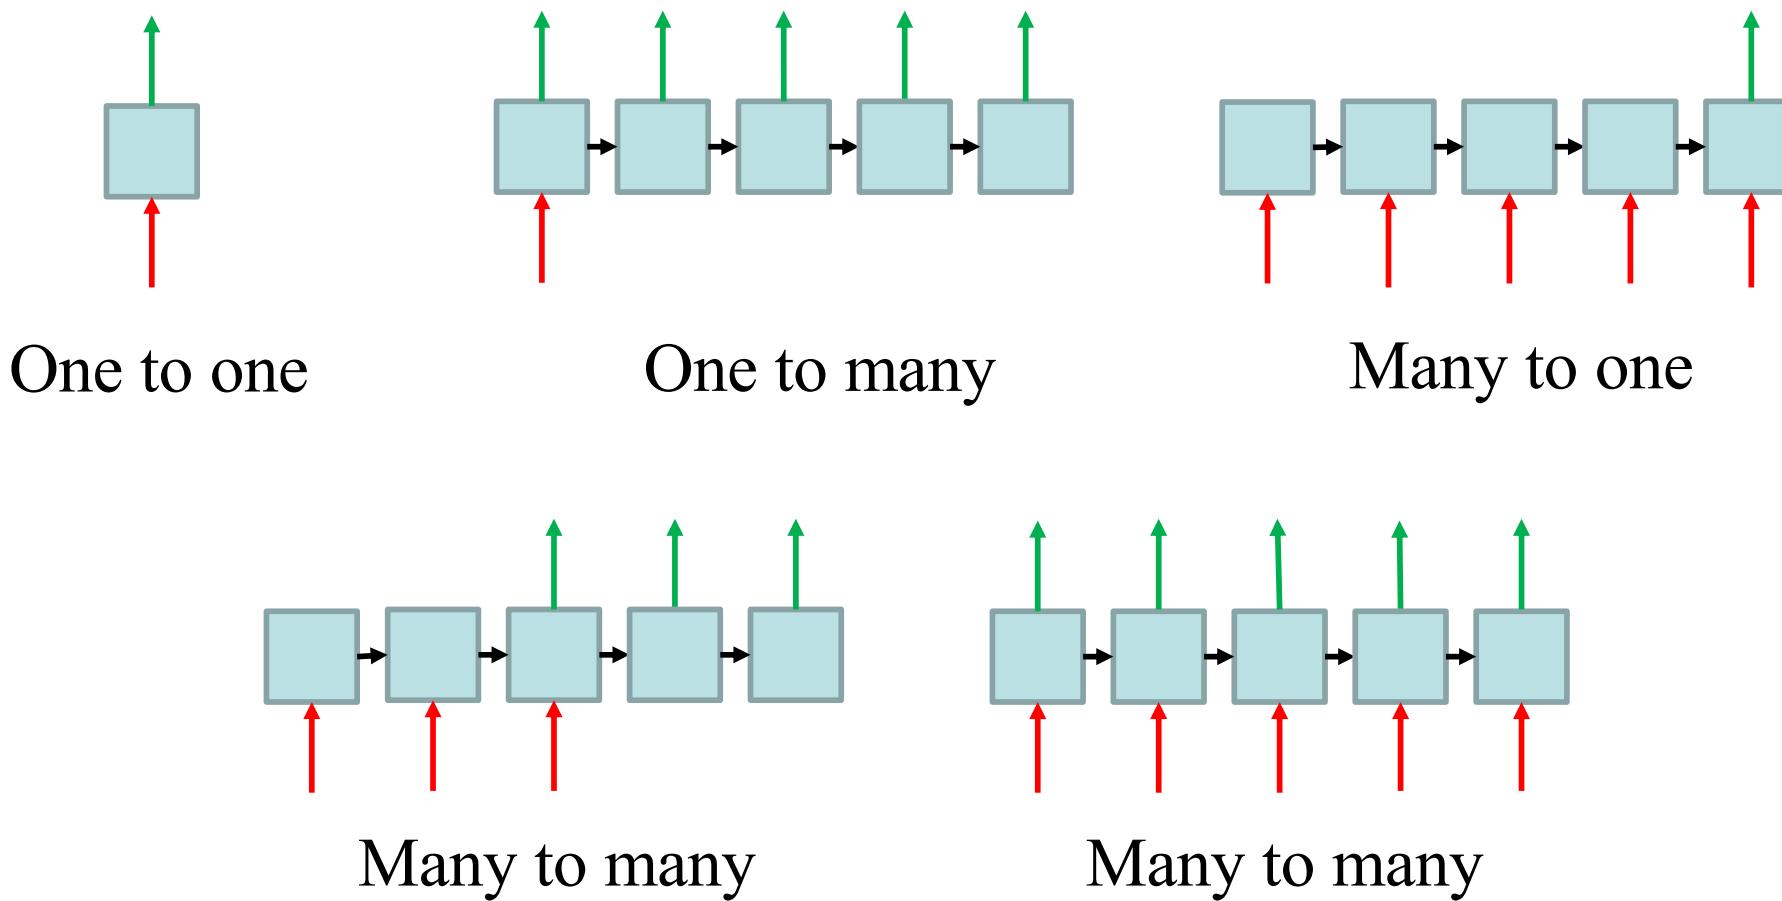
Loss function for t

→ Forward
→ Backward



Recurrent Neural Networks

◆ Flexible architecture



Recurrent Neural Networks

- ◆ Bidirectional RNN: Information from past and future states
- ◆ Forward pass:

Forward states: $h_{t-1} \rightarrow h_t \rightarrow h_{t+1}$

Backward states: $h'_{t+1} \rightarrow h'_t \rightarrow h'_{t-1}$

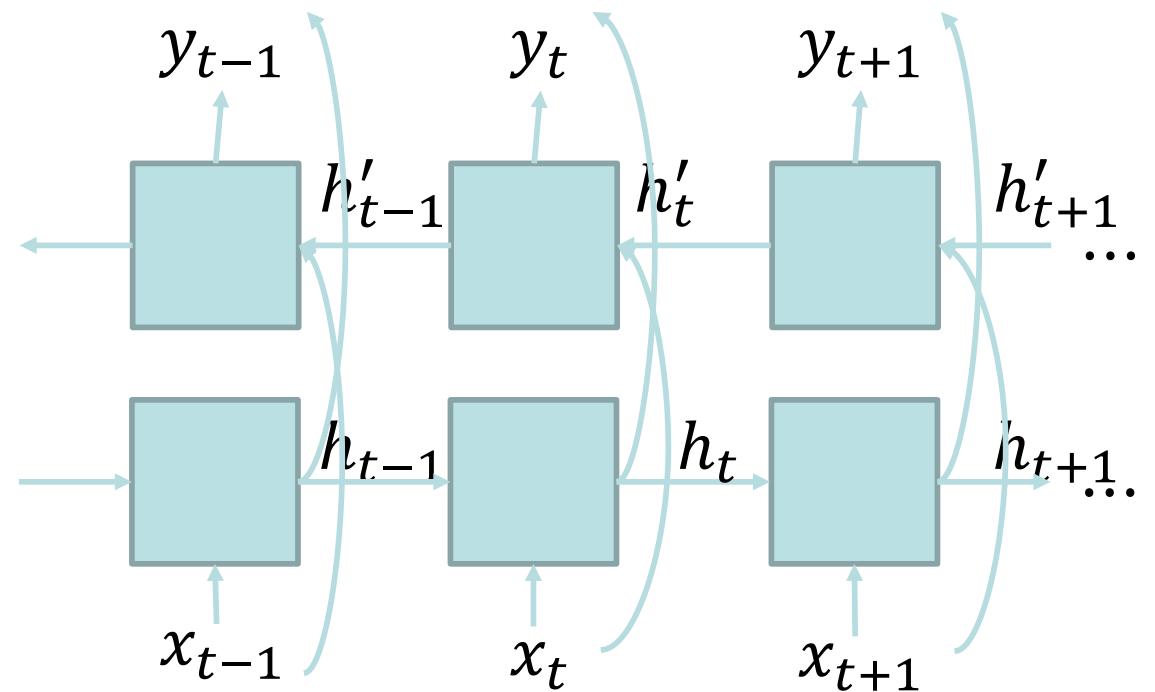
- ◆ Backward pass:

Forward states:

$h_{t+1} \rightarrow h_t \rightarrow h_{t-1}$

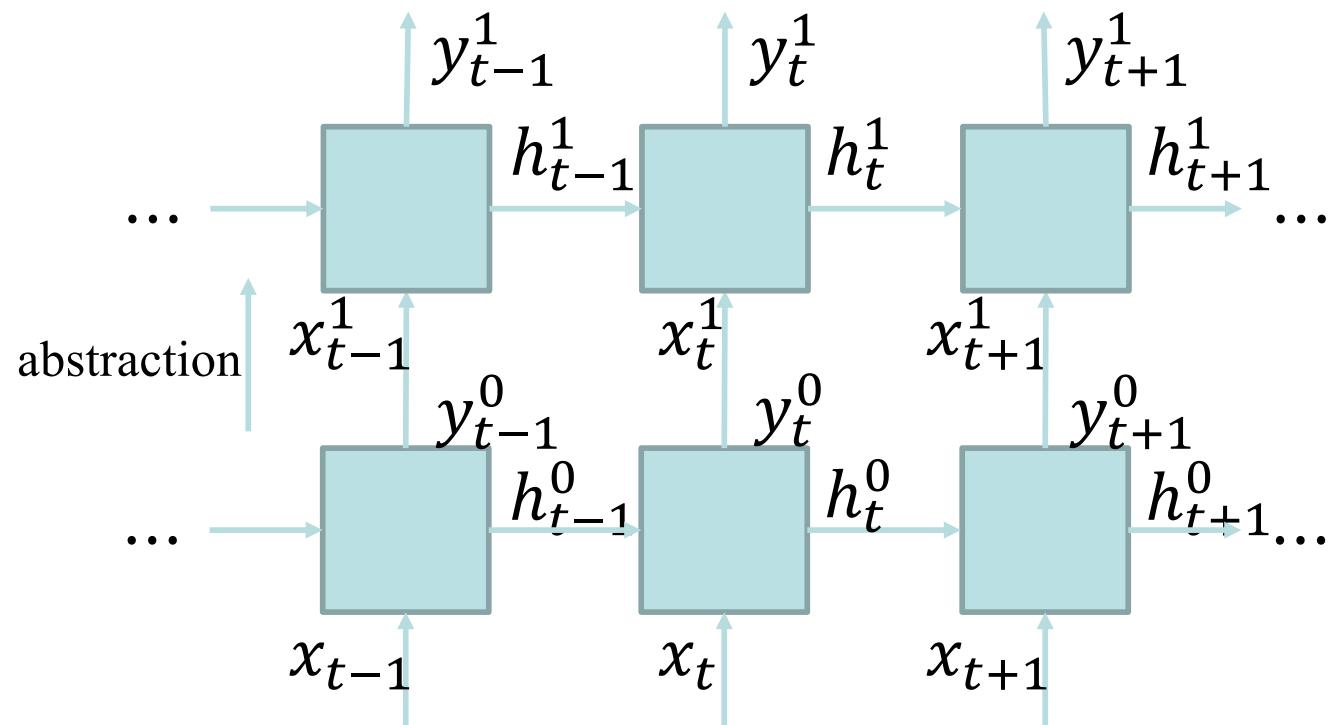
Backward states:

$h'_{t-1} \rightarrow h'_t \rightarrow h'_{t+1}$



Recurrent Neural Networks

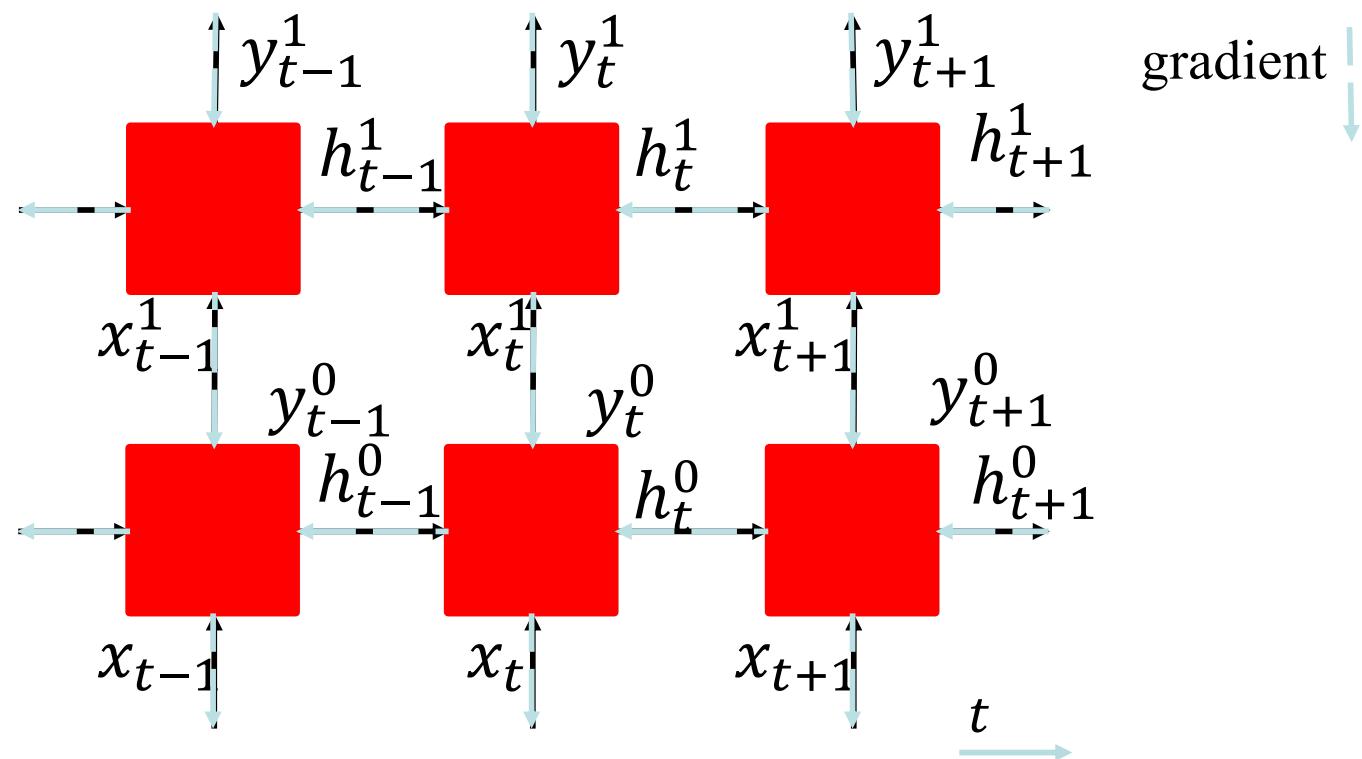
- ◆ Deep RNN: Stack hidden layers
- ◆ Higher layers for higher-layer features



Recurrent Neural Networks

◆ Backpropagation for optimization

- Forward: Along t layer by layer
- Backward: Update with gradient descent



Long Short-Term Memory (LSTM)

- ◆ **Input:** h_{t-1} , c_{t-1} and x_t
- ◆ **Forget gate:** $f_t = \sigma(x_t w_{xf} + h_{t-1} w_{hf} + b_f)$

- ◆ **Input gate:**

$$i_t = \sigma(x_t w_{xi} + h_{t-1} w_{hi} + b_i)$$

$$g_t = \tanh(x_t w_{xg} + h_{t-1} w_{hg} + b_g)$$

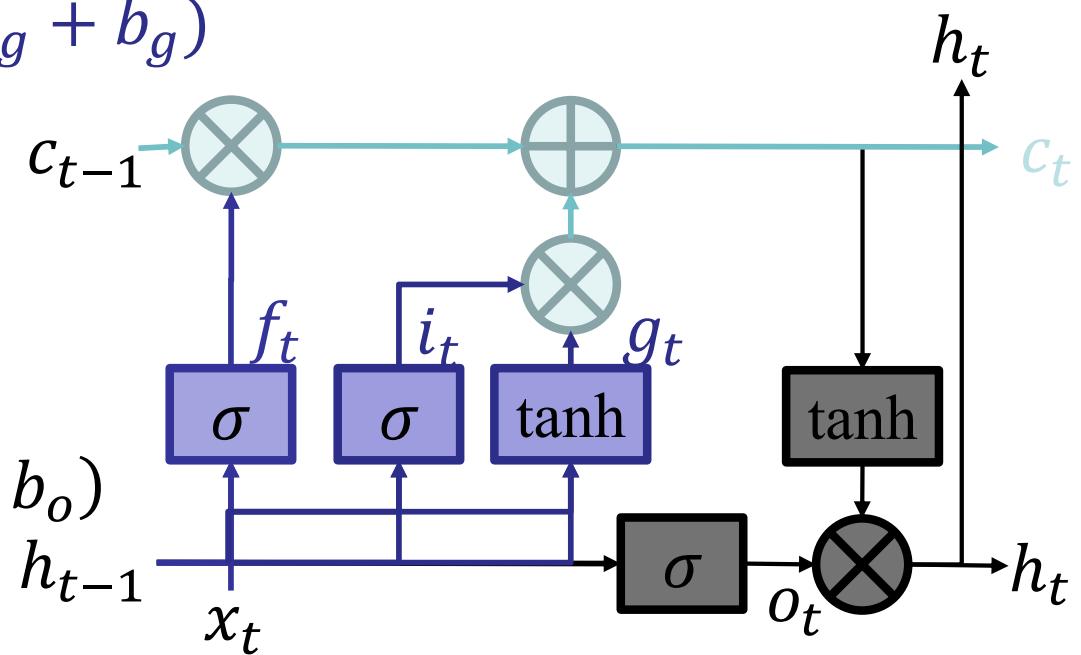
- ◆ **Current state:**

$$c_t = f_t \odot c_{t-1} + i_t \odot g_t$$

- ◆ **Output:**

$$o_t = \sigma(x_t w_{xo} + h_{t-1} w_{ho} + b_o)$$

$$h_t = o_t \odot \tanh(c_t)$$



Gated Recurrent Unit (GRU)

- ◆ Input: h_{t-1} and x_t
- ◆ **Reset gate:** $r_t = \sigma(x_t w_{xr} + h_{t-1} w_{hr} + b_r)$

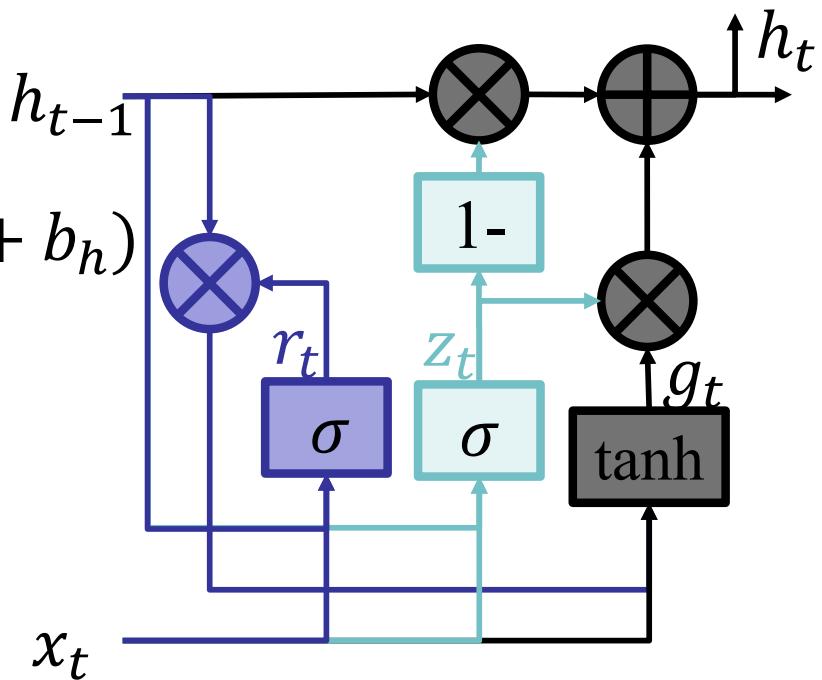
- ◆ Update gate:

$$z_t = \sigma(x_t w_{xz} + h_{t-1} w_{hz} + b_z)$$

- ◆ Output:

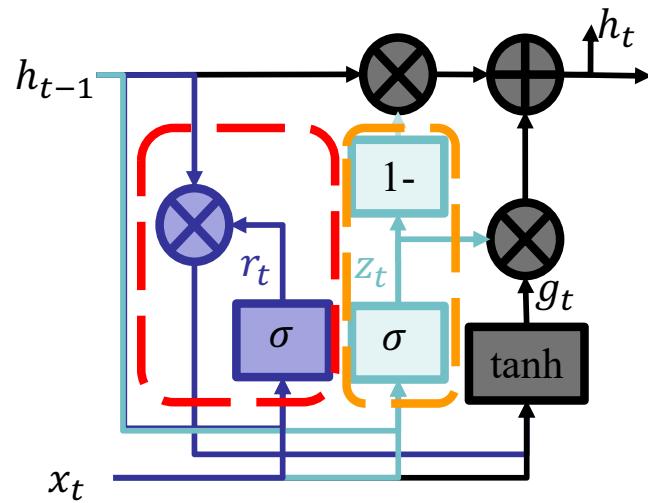
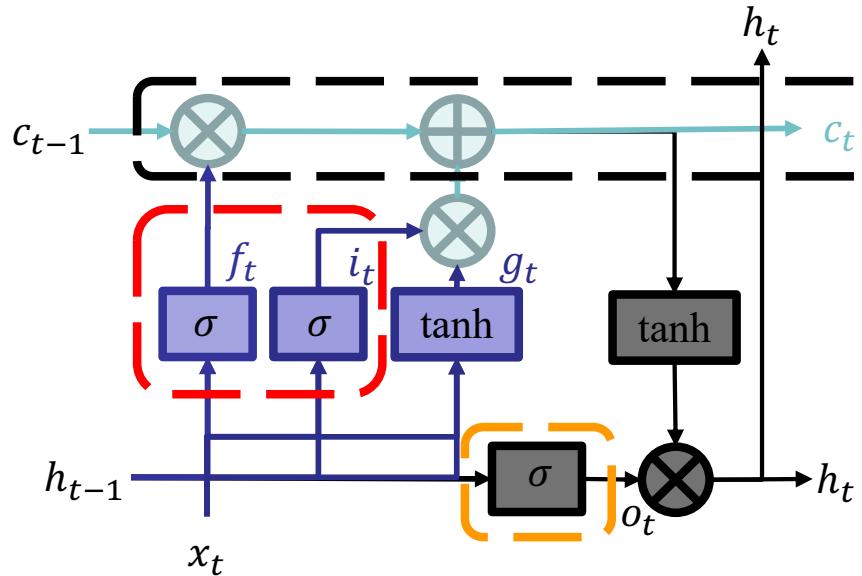
$$g_t = \tanh(x_t w_{xh} + (r_t \odot h_{t-1}) w_{hh} + b_h)$$

$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot g_t$$



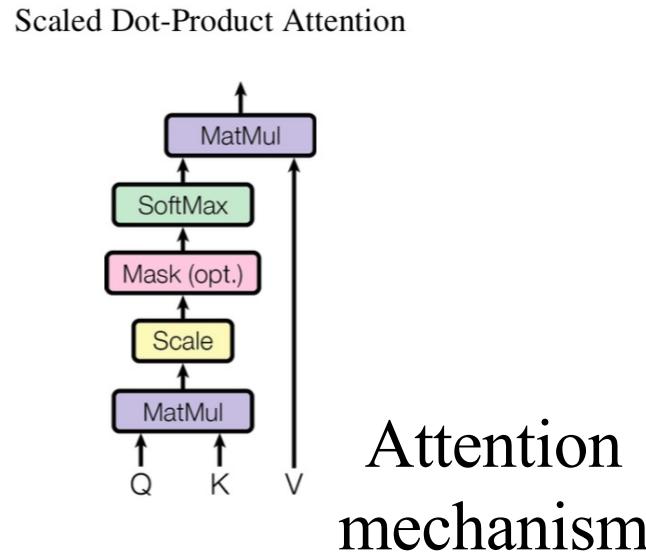
LSTM vs. GRU

- ◆ GRU is more compressed in comparison to LSTM:
 - GRU merges cell state and hidden state;
 - GRU combines forget gate and input into update gate
- ◆ GRU directly uses h_t to retain information from previous cell or jump start with new information



Transformer

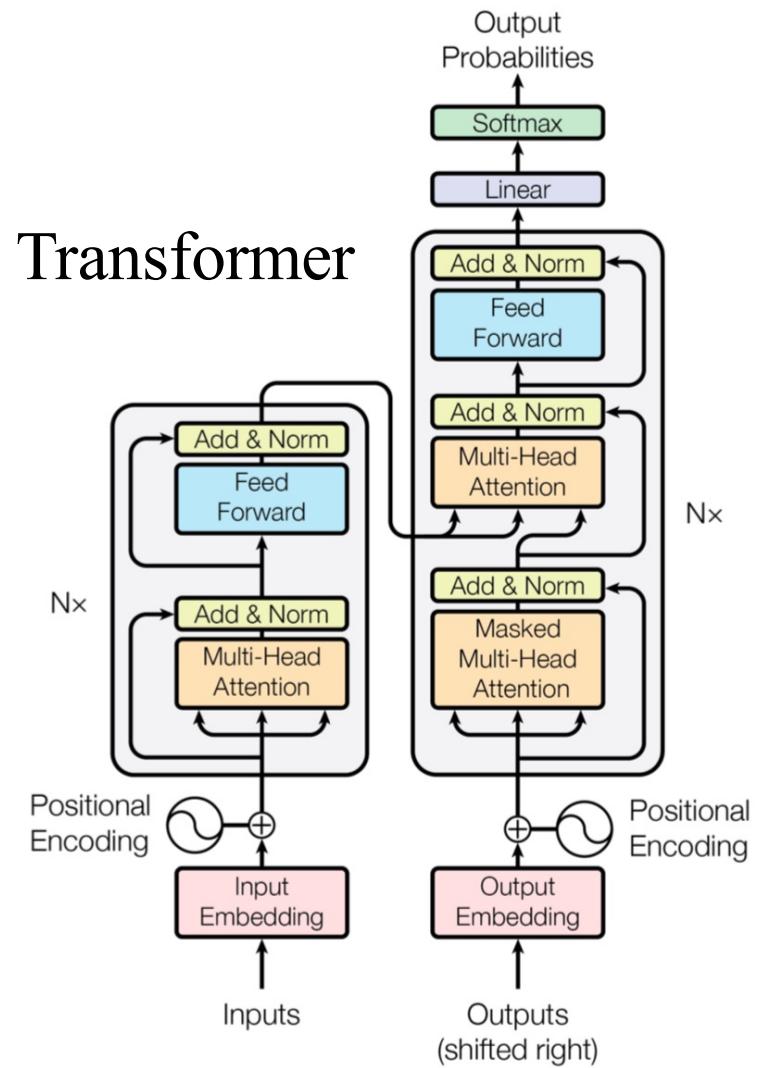
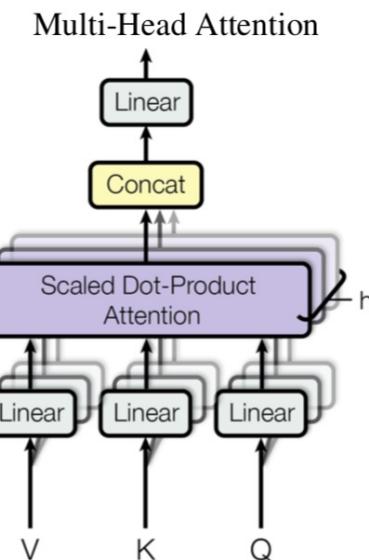
- ◆ Feed forward network + Self attention



$$\begin{aligned} Q^T &= X^T W_q^T \\ K^T &= X^T W_k^T \\ V^T &= X^T W_v^T \end{aligned}$$

$$S^\tau = \text{softmax}\left(\frac{Q^\tau(K^\tau)^T}{\sqrt{d_k^\tau}}\right)$$

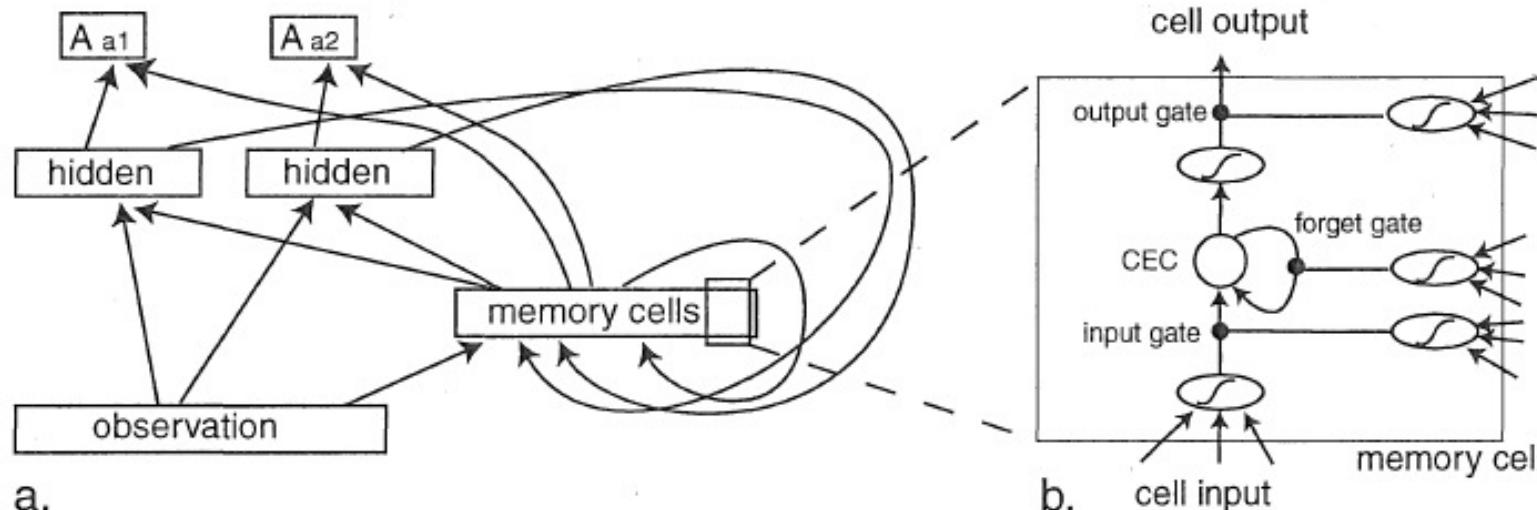
$$M^\tau = S^\tau V^\tau$$



Example: Non-Markovian RL

- ◆ **RL-LSTM:** Model-free RL with LSTM for non-Markovian tasks with long-term dependencies
- ◆ Directly approximate the value function of advantage learning and LSTM to backpropagate prediction error

$$E^{TD}(t) = V(s(t)) + \frac{r(t) + \gamma V(s(t+1)) - V(s(t))}{\kappa} - A(s(t), a(t))$$



Bakker, Bram. "Reinforcement learning with long short-term memory." In *Advances in neural information processing systems*, pp. 1475-1482. 2002.

Generative Models

Generative vs. Discriminative

◆ Discriminative models

- Predict a label or category given the features of an instance of data
- Learn the **boundary** between classes

◆ Generative models:

- Predict features given a certain label
- Model the **distribution** of individual classes

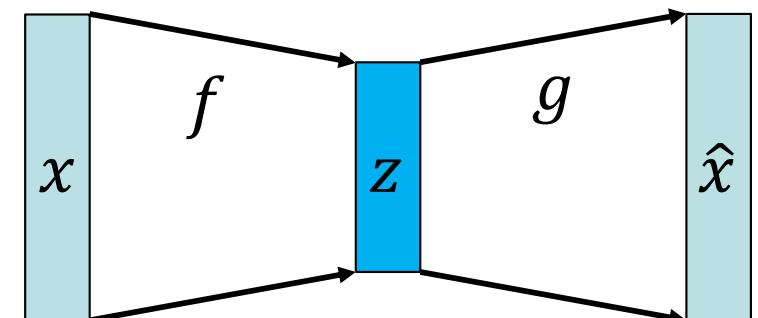
Autoencoders

◆ Neural networks trained to copy its input to its output

- **Encoder:** map the input x to a hidden representation z
- **Decoder:** map the hidden representation z to the output \hat{x}
- **Bottleneck layer** with lower dimension
- Loss function for training: $L(x, \hat{x}) = L(x, g(f(x)))$

◆ Regularization

- ◆ Sparsity: Sparse Autoencoder
- ◆ Robustness: Denoising Autoencoder



◆ Dimensionality reduction & PCA

Encoder
 $z = f(x)$

Decoder
 $\hat{x} = g(z)$

Sparse Autoencoders

- ◆ Sparsity penalty on Bottleneck layer

$$J_s(W, b) = J(W, b) + \beta \sum_{j=1}^S KL(\rho || \hat{\rho}_j)$$

- ◆ KL divergence for penalty on hidden representations

$$KL(\rho || \hat{\rho}_j) = \rho \log \frac{\rho}{\hat{\rho}_j} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_j}$$



Expected activation



Average activation for j -th hidden unit

- ◆ For j -th hidden representation

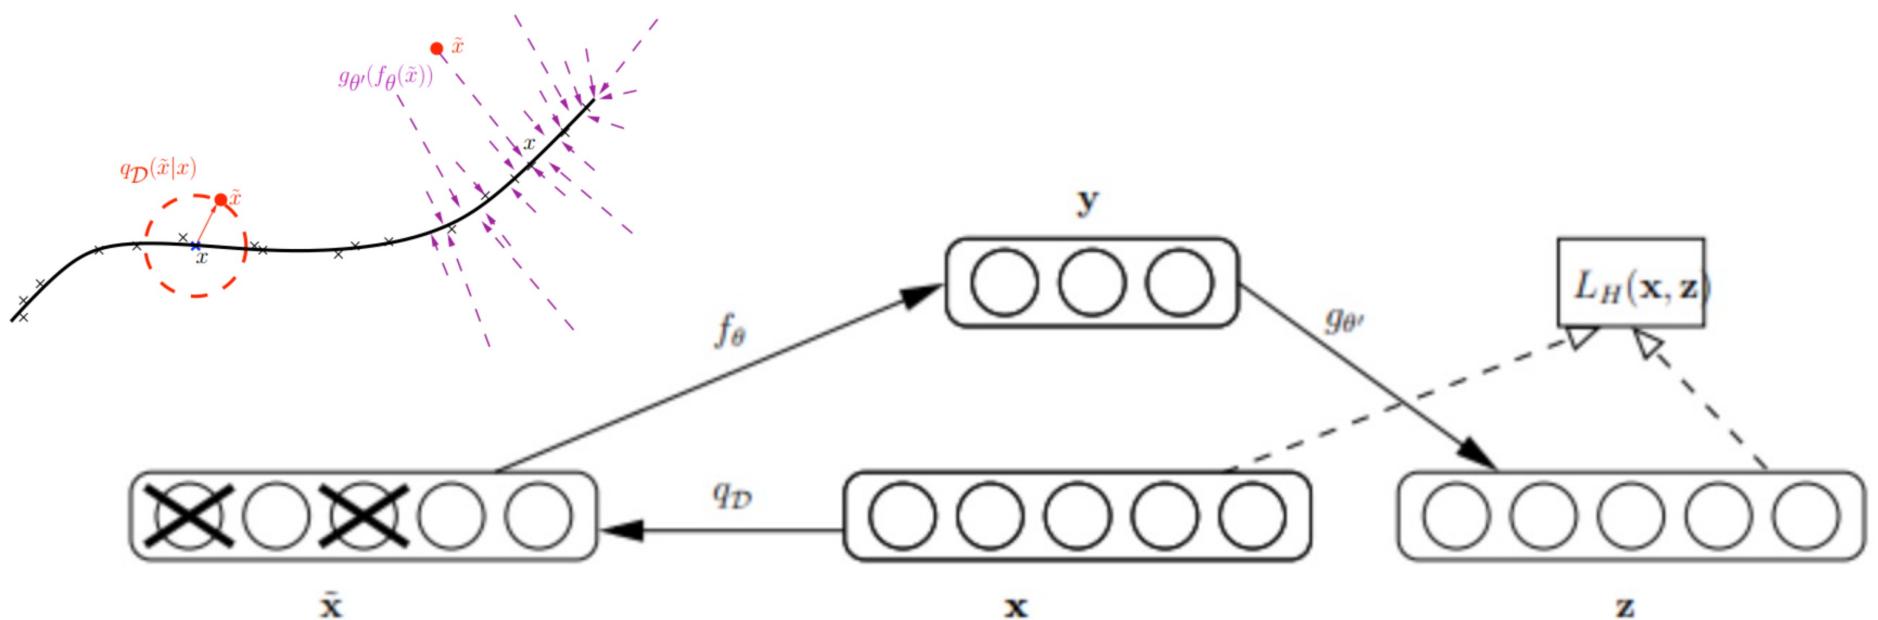
$$\hat{\rho}_j = \frac{1}{m} \sum_{i=1}^m [a_j^{(2)}(x_i)]$$



Activation for specific $input x$

Denoising Autoencoders

- ◆ Noise perturbation on input x to **eliminate partial components**
- ◆ Minimize loss $L_H(x, z)$ between x and z
- ◆ Relation to Dropout (on hidden units)



Variational Autoencoders (VAEs)

- ◆ Autoencoder with regularization on distribution for encoder/decoder to organize the latent space for generation

- ◆ Encode the input x as distribution $p_\theta(z|x)$ over the latent space
- ◆ Decode the sample from the latent space using the distribution
- ◆ Backpropagate the reconstruction error for training

$$L(\theta, \phi) = \sum_i -\mathbb{E}_{z \sim p_\theta(z|x_i)} [\log q_\phi(x_i|z)] + \mathbb{D}_{\text{KL}}(p_\theta(z|x_i) || p(z))$$

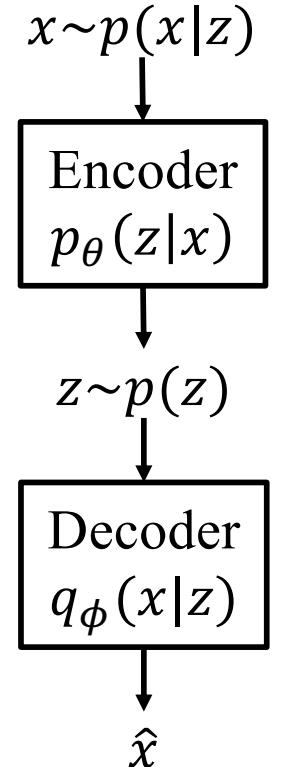
Reconstruction loss

Regularization

- ◆ Equivalent to ELBO in variational inference

- ◆ Reparametrization trick

$$\varepsilon \sim \mathcal{N}(0,1) \xrightarrow{z = \mu + \varepsilon \cdot \sigma} z \sim \mathcal{N}(\mu, \sigma^2) \xrightarrow{\mu = f(x), \log \sigma = g(x)} \hat{x}$$

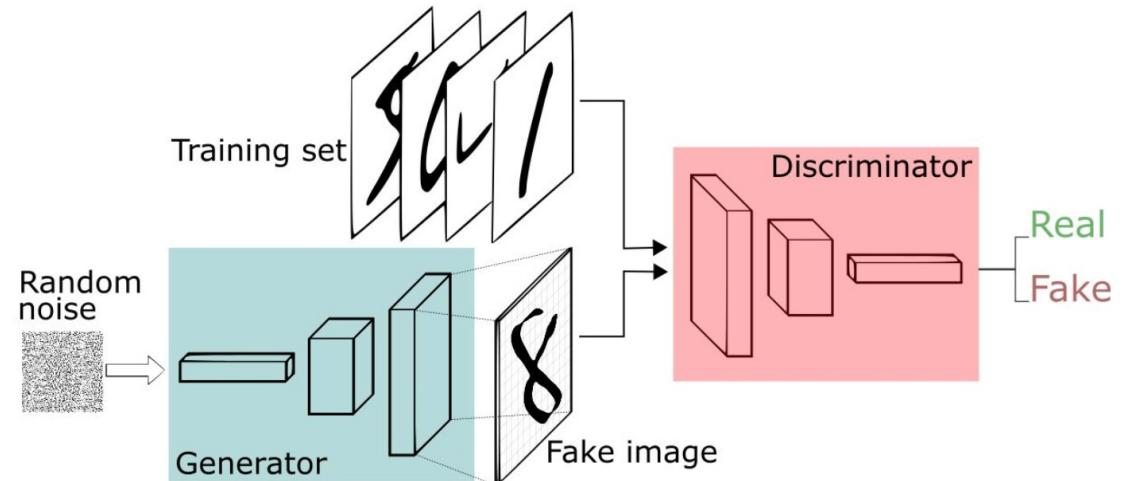


Generative Adversarial Networks

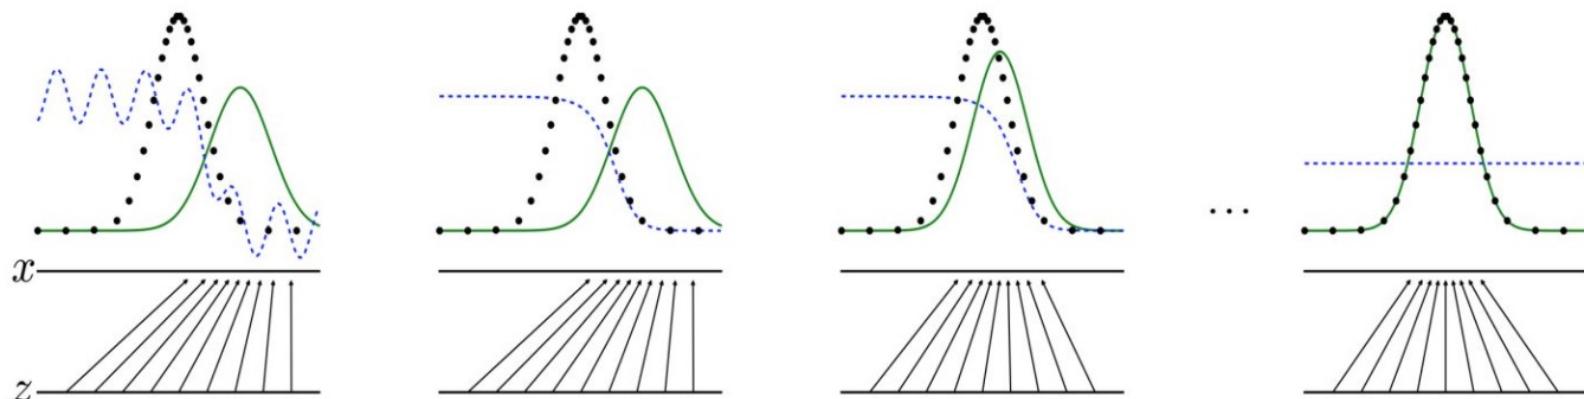
◆ Generator

◆ Discriminator

◆ Update Alternately



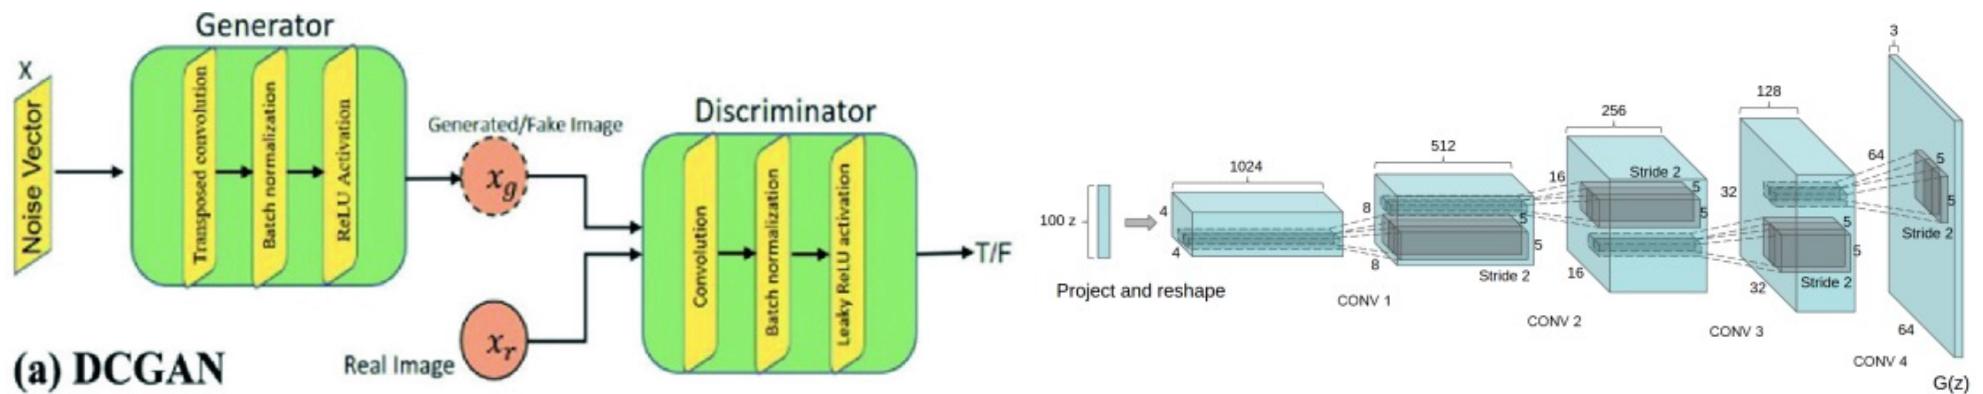
$$\min_G \max_D L(G, D) = \mathbb{E}_{x \sim p_X(x)}[\log D(x)] + \mathbb{E}_{z \sim p_Z(z)}[\log(1 - D(G(z)))]$$



Deep Convolutional GAN

◆ DCGAN: Combining CNN and GAN

- ◆ Replace pooling layers with **fractional-stride convolutions**
- ◆ *Batch Normalization* and *LeakyReLU* activation
- ◆ Fully convolution network



Radford, A., Metz, L., & Chintala, S. (2016). Unsupervised representation learning with deep convolutional generative adversarial networks. In 4th International Conference on Learning Representations ICLR 2016—Conference Track Proceedings (pp. 1–16).

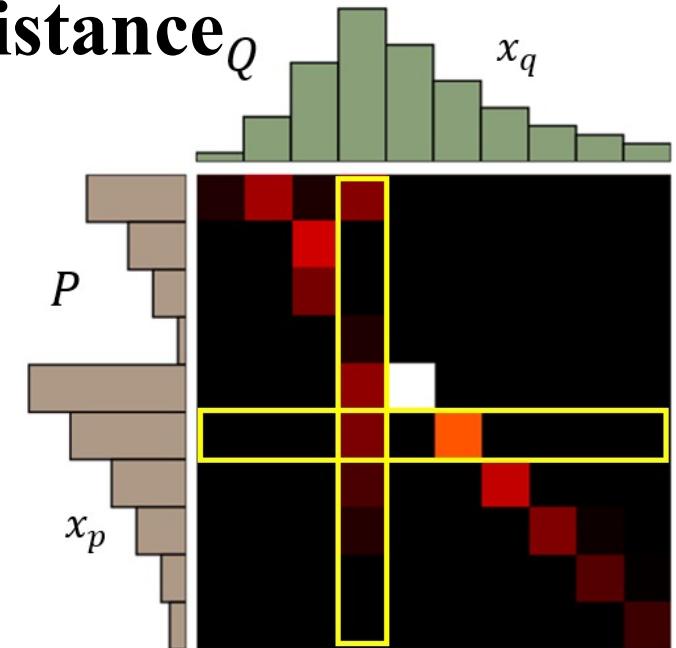
Wasserstein GAN

◆ WGAN: GAN using Wasserstein distance

- ◆ Earth mover & moving plan

Average distance of a plan γ

$$B(\gamma) = \sum_{x_p, x_q} \gamma(x_p, x_q) \|x_p - x_q\|$$



$$\text{Wasserstein distance } W(P, Q) = \min_{\gamma \in \Pi} B(\gamma)$$

- ◆ Optimization

$$V(G, D) = \max_{D \in \mathcal{C}} \mathbb{E}_{x \sim p_{\text{data}}} [D(x)] + \mathbb{E}_{x \sim p_G} [D(x)]$$

$$\mathcal{C} = \{D \mid \|D(x_1) - D(x_2)\| \leq \|x_1 - x_2\|\}$$

1-Lipschitz

- ◆ WGAN-GP & CPGAN

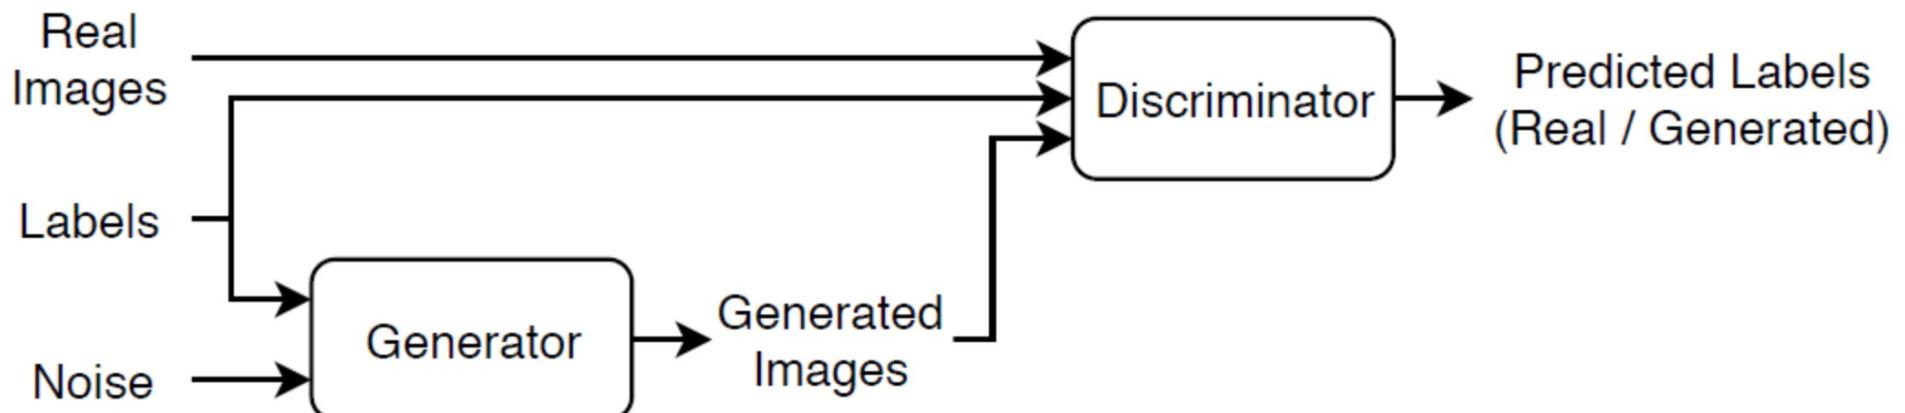
Conditonal GAN

◆ cGAN: GAN using labels during training

- Generator: generate data with the same structure as the training data observations corresponding to the same label.
- Discriminator: classify the observations as "real" or "generated", given training and generated data, this network

$$\min_G \max_D L(G, D)$$

$$= \mathbb{E}_{x \sim p_X(x)} [\log D(x|y)] + \mathbb{E}_{z \sim p_Z(z)} [\log(1 - D(G(z|y)))]$$

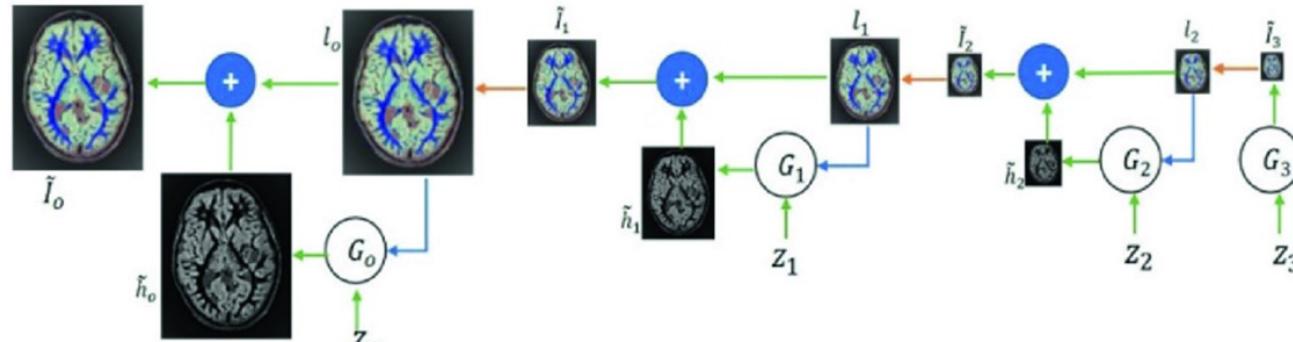


LAPGAN

◆ LAPGAN: Successive sampling for full resolution

- ◆ A set of generative models $\{G_0, G_1, \dots, G_K\}$ at each level k of the Laplacian Pyramid $L(I)$ to capture the coefficients h_k for input image I
- ◆ h_k is calculated by difference of adjacent levels in Gaussian Pyramid $G_k(I)$ and up-sampled value $U(G_{k+1}(I))$

$$h_k = L_k(I) = G_k(I) - U(G_{k+1}(I)) = I_k - U(I_{k+1})$$

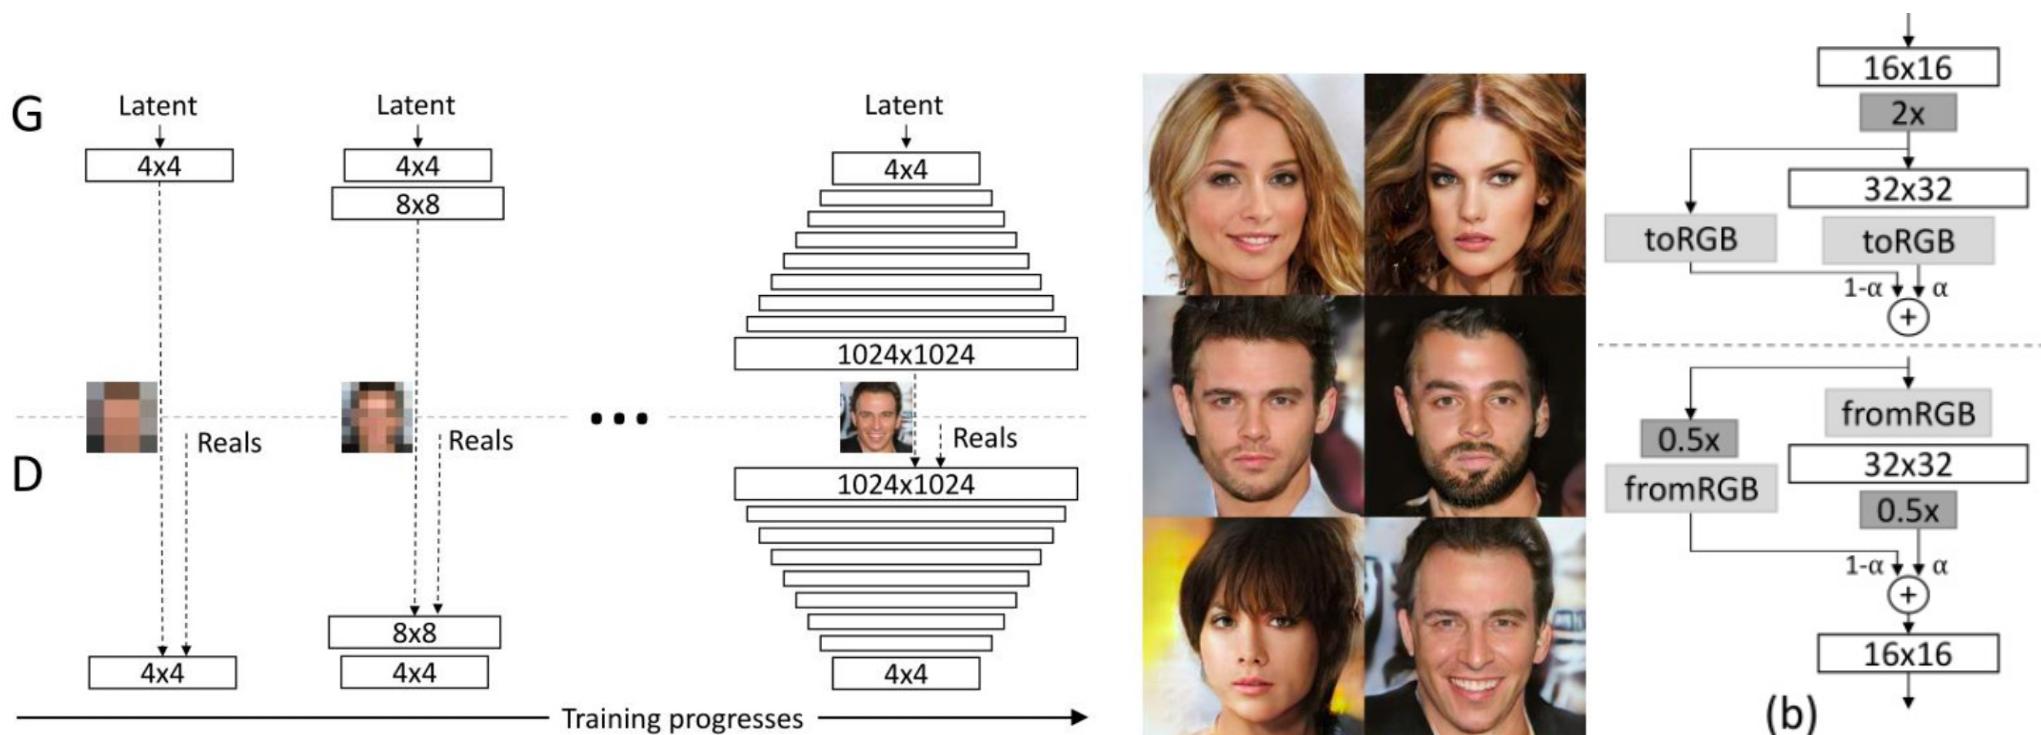


(b) LAPGAN

ProGAN

◆ ProGAN: Progressive Growing GAN

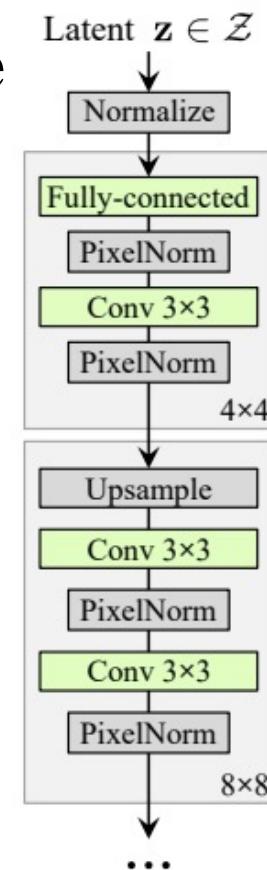
- ◆ Start from low resolution and refine details
- ◆ Usage of Skip connection



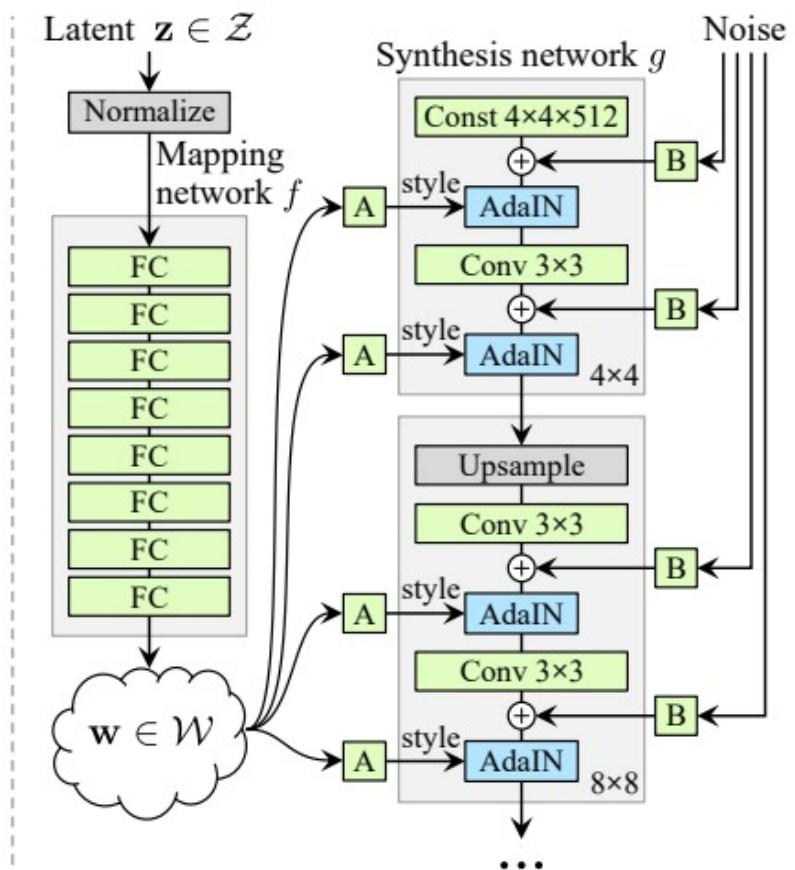
StyleGAN

◆ StyleGAN: Regulating outputs

- ◆ Style-based generator:
Sequentially produce the simulated image
- ◆ Transform the input of each level individually without altering other levels



(a) Traditional



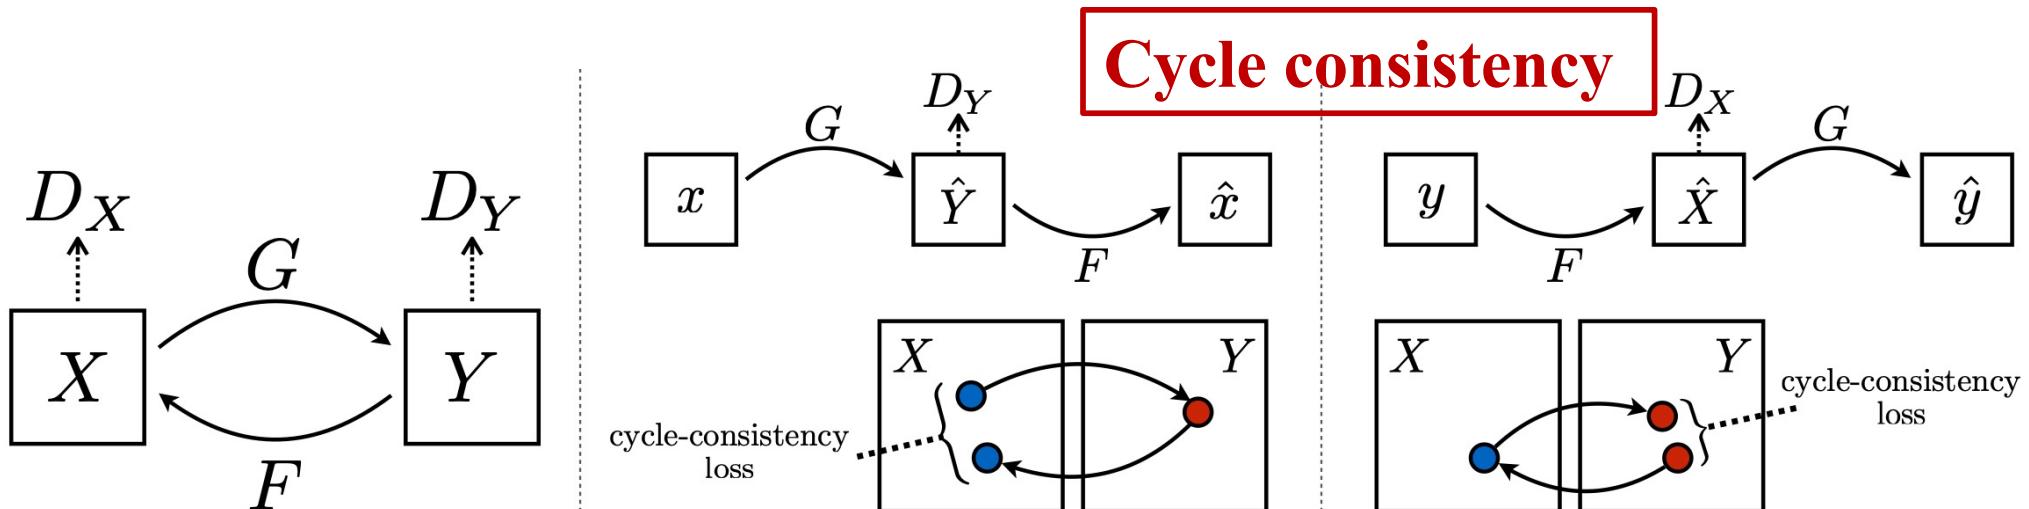
(b) Style-based generator

CycleGAN

- ◆ Absence of paired training data. Learn mapping function between two domains X and Y
- ◆ Two adversarial discriminators D_X and D_Y

$$\min_{G,F} \max_{D_X,D_Y} \mathcal{L} = L_{GAN}(G, D_Y, X, Y) + L_{GAN}(F, D_X, Y, X) + L_{cyc}(G, F)$$

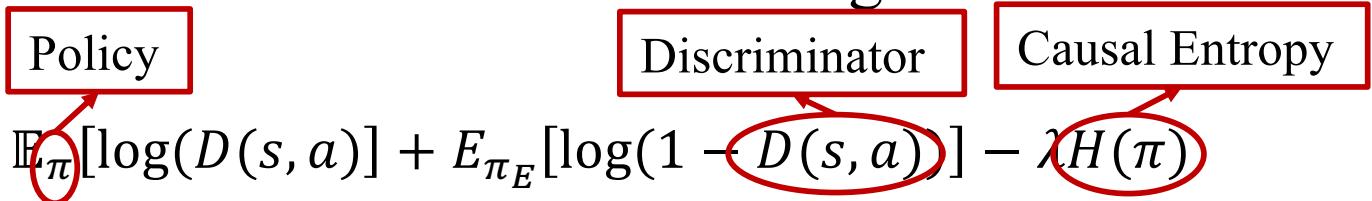
$$L_{cyc}(G, F) = \mathbb{E}_{x \sim p(x)} [\|F(G(x)) - x\|_1] + \mathbb{E}_{y \sim p(y)} [\|G(F(y)) - y\|_1]$$



Jun-Yan Zhu*, Taesung Park*, Phillip Isola, and Alexei A. Efros. "Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks", in IEEE International Conference on Computer Vision (ICCV), 2017.

Example: Generative Adversarial Imitation Learning

- ◆ Directly learn a policy from expert trajectories without inverse RL
- ◆ Fit distributions of states and actions with generative adversarial training



Algorithm 1 Generative adversarial imitation learning

- 1: **Input:** Expert trajectories $\tau_E \sim \pi_E$, initial policy and discriminator parameters θ_0, w_0
- 2: **for** $i = 0, 1, 2, \dots$ **do**
- 3: Sample trajectories $\tau_i \sim \pi_{\theta_i}$
- 4: Update the discriminator parameters from w_i to w_{i+1} with the gradient

$$\hat{\mathbb{E}}_{\tau_i}[\nabla_w \log(D_w(s, a))] + \hat{\mathbb{E}}_{\tau_E}[\nabla_w \log(1 - D_w(s, a))] \quad (17)$$

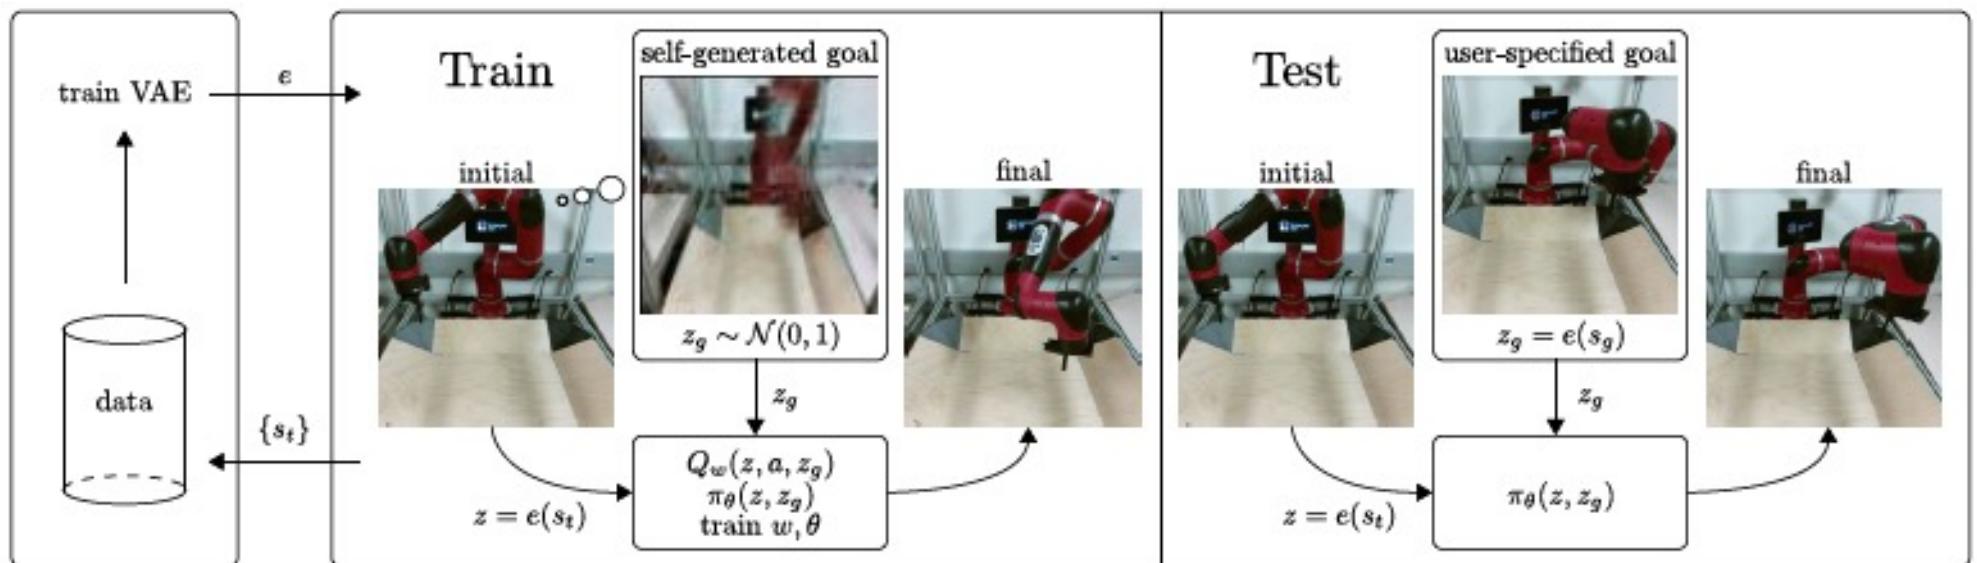
- 5: Take a policy step from θ_i to θ_{i+1} , using the TRPO rule with cost function $\log(D_{w_{i+1}}(s, a))$. Specifically, take a KL-constrained natural gradient step with

$$\begin{aligned} & \hat{\mathbb{E}}_{\tau_i} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q(s, a)] - \lambda \nabla_{\theta} H(\pi_{\theta}), \\ & \text{where } Q(\bar{s}, \bar{a}) = \hat{\mathbb{E}}_{\tau_i} [\log(D_{w_{i+1}}(s, a)) \mid s_0 = \bar{s}, a_0 = \bar{a}] \end{aligned} \quad (18)$$

- 6: end for**

Example: Visual RL

- ◆ Combine unsupervised representation learning and RL
- ◆ Train a VAE using data generated by exploration policy
 - Sample goals to train the policy
 - Embed the observation into a latent space
 - Compute distances in the latent space



Nair, Ashvin V., Vitchyr Pong, Murtaza Dalal, Shikhar Bahl, Steven Lin, and Sergey Levine. "Visual reinforcement learning with imagined goals." In *Advances in Neural Information Processing Systems*, pp. 9191-9200. 2018.

Q & A



Many Thanks